Unit 1 Lecture 4: Exploratory data analysis

September 9, 2021

Welcome back to STAT 471! We are now in Unit 1 Lecture 4:

Unit 1: Intro to modern data mining	**Lecture 1:** Intro to modern data mining
Unit 2: Tuning predictive models	**Lecture 2:** Linear regression
Unit 3: Regression-based methods	**Lecture 3:** Data wrangling
Unit 4: Tree-based methods	**Lecture 4:** Exploratory data analysis
Unit 5: Deep learning	**Lecture 5:** Unit review and quiz in class
	Homework 1 due the following Sunday.

This lecture is about *exploratory data analysis*, which involves data transformation and visualization. It draws on Chapters 3, 5, and 7 from the excellent R for Data Science book (direct quotations are presented using block quotes).

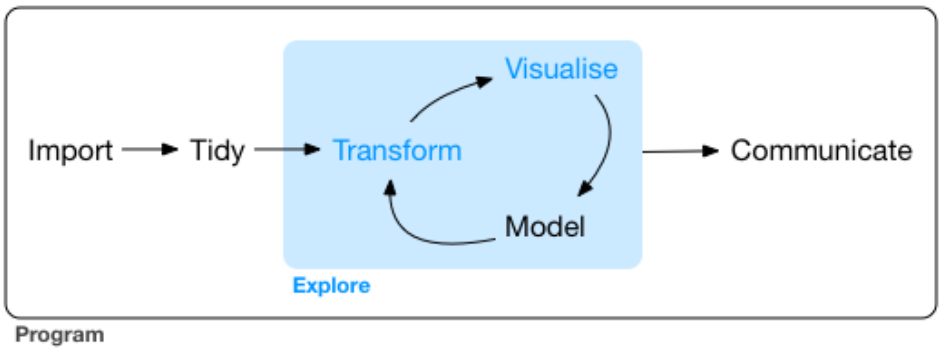

Figure 1: Image source: R4DS Chapter 2.

As usual, let's load the `tidyverse`:

```
library(tidyverse)
```

1 Data visualization

> R has several systems for making graphs, but `ggplot2` [one of the core members of the `tidyverse`] is one of the most elegant and most versatile. `ggplot2` implements the grammar of graphics, a coherent system for describing and building graphs. With `ggplot2`, you can do more faster by learning one system and applying it in many places.

1.1 `ggplot` basics

Let's recall the `mpg` data frame from last lecture:

```
mpg
```

```
## # A tibble: 234 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 audi         a4           1.8  1999     4 auto~ f        18    29 p     comp~
##  2 audi         a4           1.8  1999     4 manu~ f        21    29 p     comp~
##  3 audi         a4           2    2008     4 manu~ f        20    31 p     comp~
##  4 audi         a4           2    2008     4 auto~ f        21    30 p     comp~
##  5 audi         a4           2.8  1999     6 auto~ f        16    26 p     comp~
##  6 audi         a4           2.8  1999     6 manu~ f        18    26 p     comp~
##  7 audi         a4           3.1  2008     6 auto~ f        18    27 p     comp~
##  8 audi         a4 quattro   1.8  1999     4 manu~ 4        18    26 p     comp~
##  9 audi         a4 quattro   1.8  1999     4 auto~ 4        16    25 p     comp~
## 10 audi         a4 quattro   2    2008     4 manu~ 4        20    28 p     comp~
## # ... with 224 more rows
```

Let's plot the relationship between `displ` (a car's engine size in liters) and `hwy` (a car's fuel efficiency on the highway, in miles per gallon).

```
mpg %>%                          # pipe in the data
  ggplot() +                     # create an empty ggplot
  geom_point(mapping =           # add scatter plot
               aes(x = displ,    # x axis location of points
                   y = hwy))     # y axis location of points
```

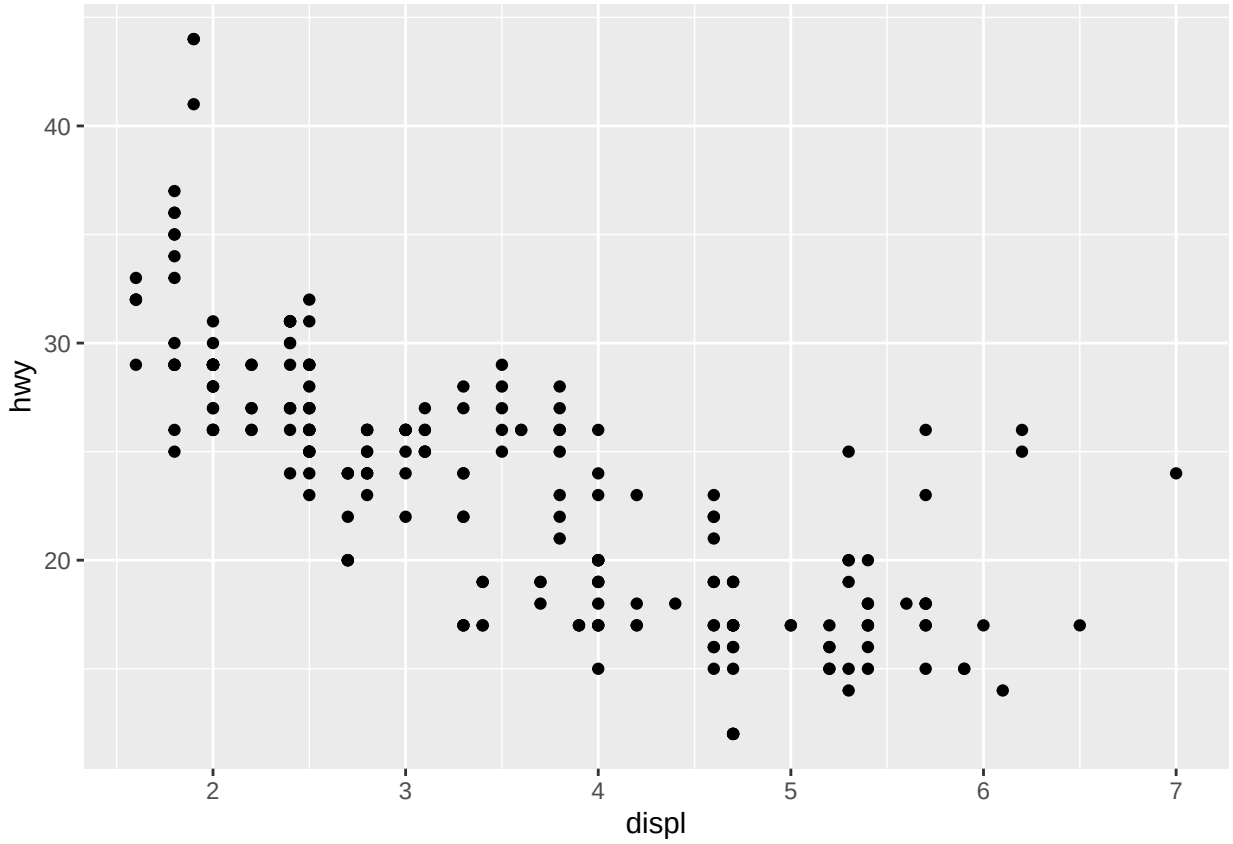

An *aesthetic* is a visual property of the objects in your plot. We create a plot by *mapping* variables in our tibble to aesthetics of the plot. In the above case, we `displ` is mapped to `x` (the horizontal axis position) and

`hwy` is mapped to `y` (the vertical axis position). A *geom* function adds a specific representation of the data to the plot (scatter plot, box plot, etc). In the above case, we used `geom_point` to create a scatter plot. A plot can have multiple geoms and multiple aesthetics. The plot above contains only one *panel*, but multi-panel plots can be created using *faceting*.

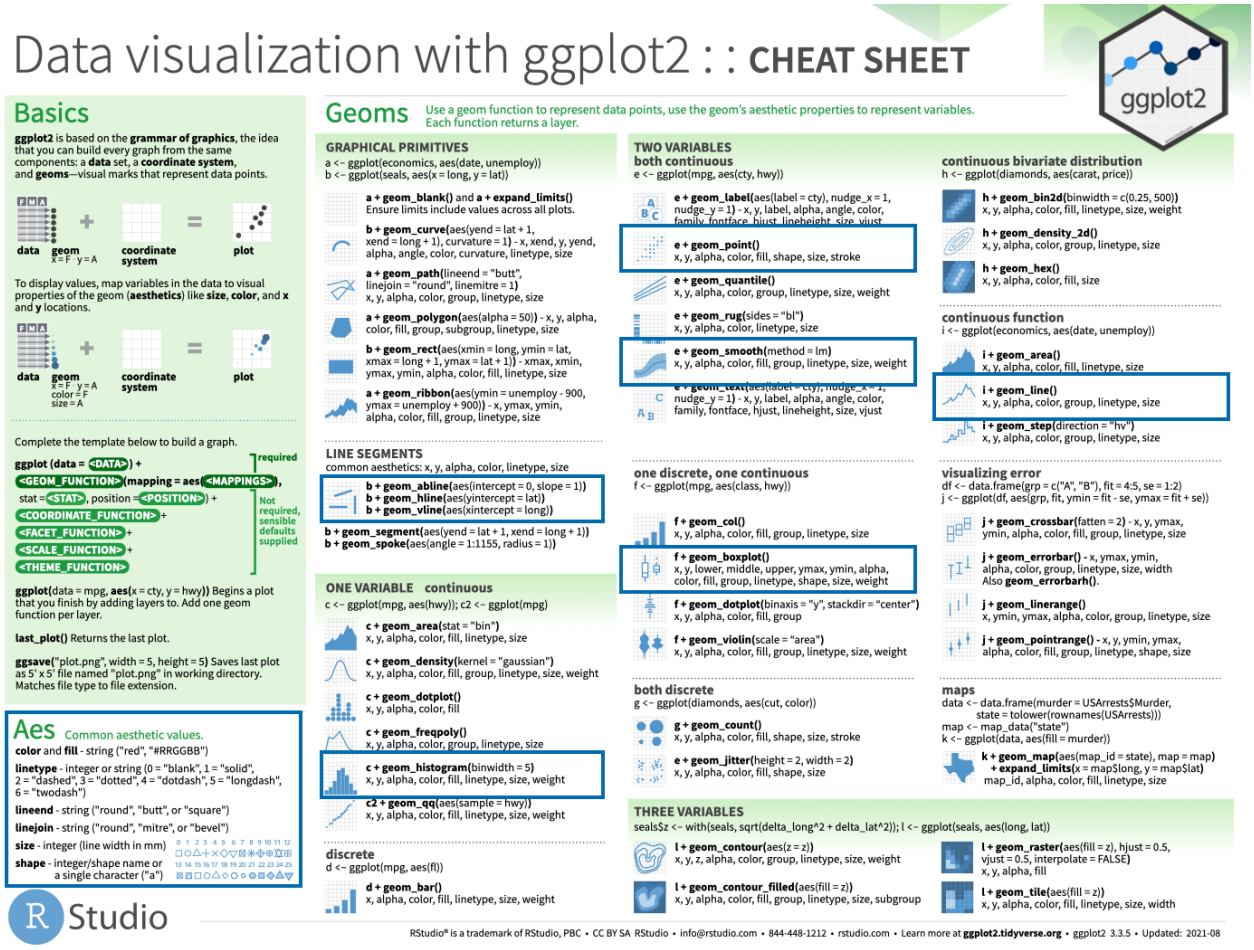

Figure 2: Image source: https://www.rstudio.com/resources/cheatsheets/

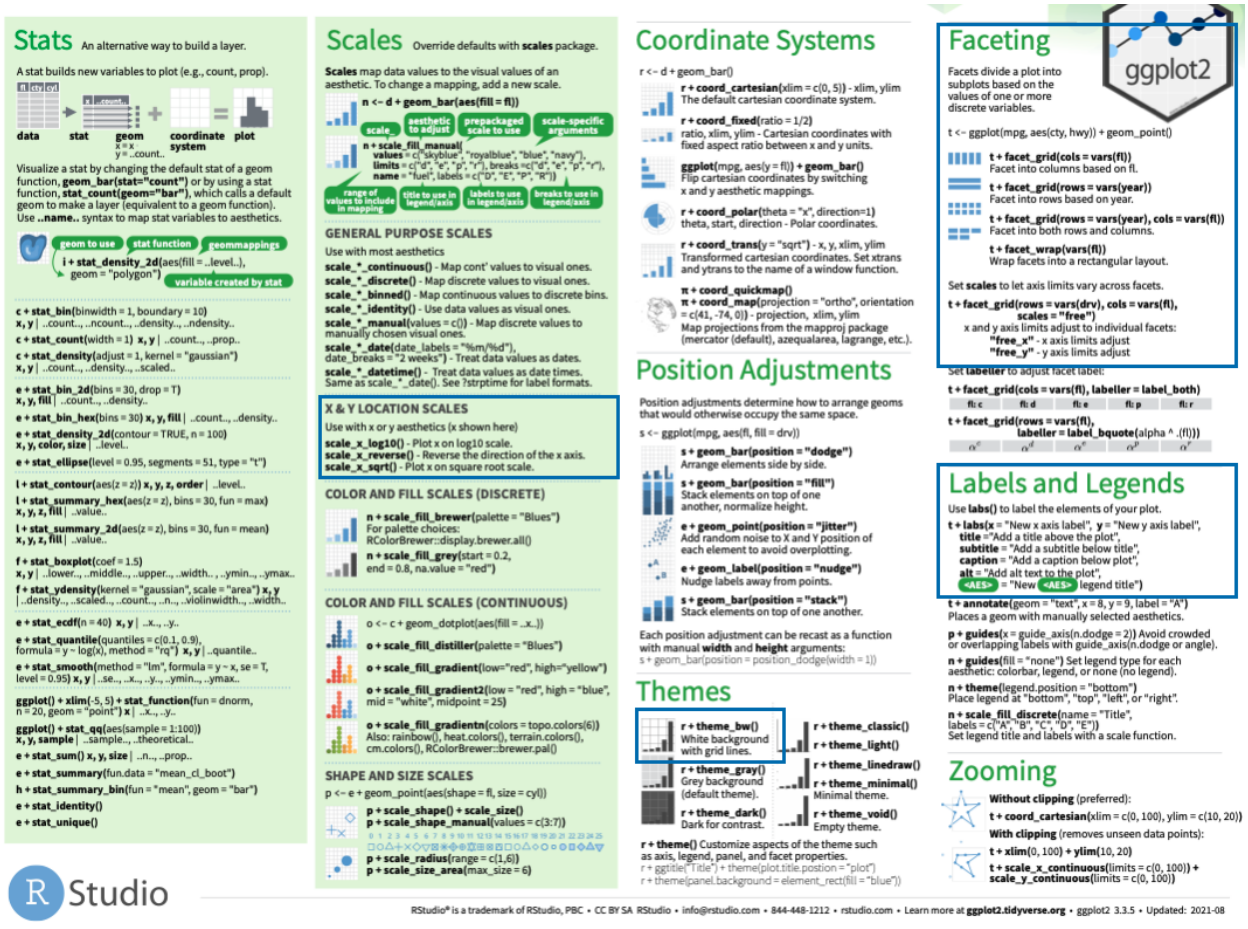

Figure 3: Image source: https://www.rstudio.com/resources/cheatsheets/

1.2 Aesthetics

Let's see some examples of different aesthetics we can add to the above scatter plot.

Adding a color aesthetic:

```
mpg %>%                          # pipe in the data
  ggplot() +                     # create an empty ggplot
  geom_point(mapping =           # add scatter plot
               aes(x = displ,    # x axis location of points
                   y = hwy,      # y axis location of points
                   color = class)) # color of points
```

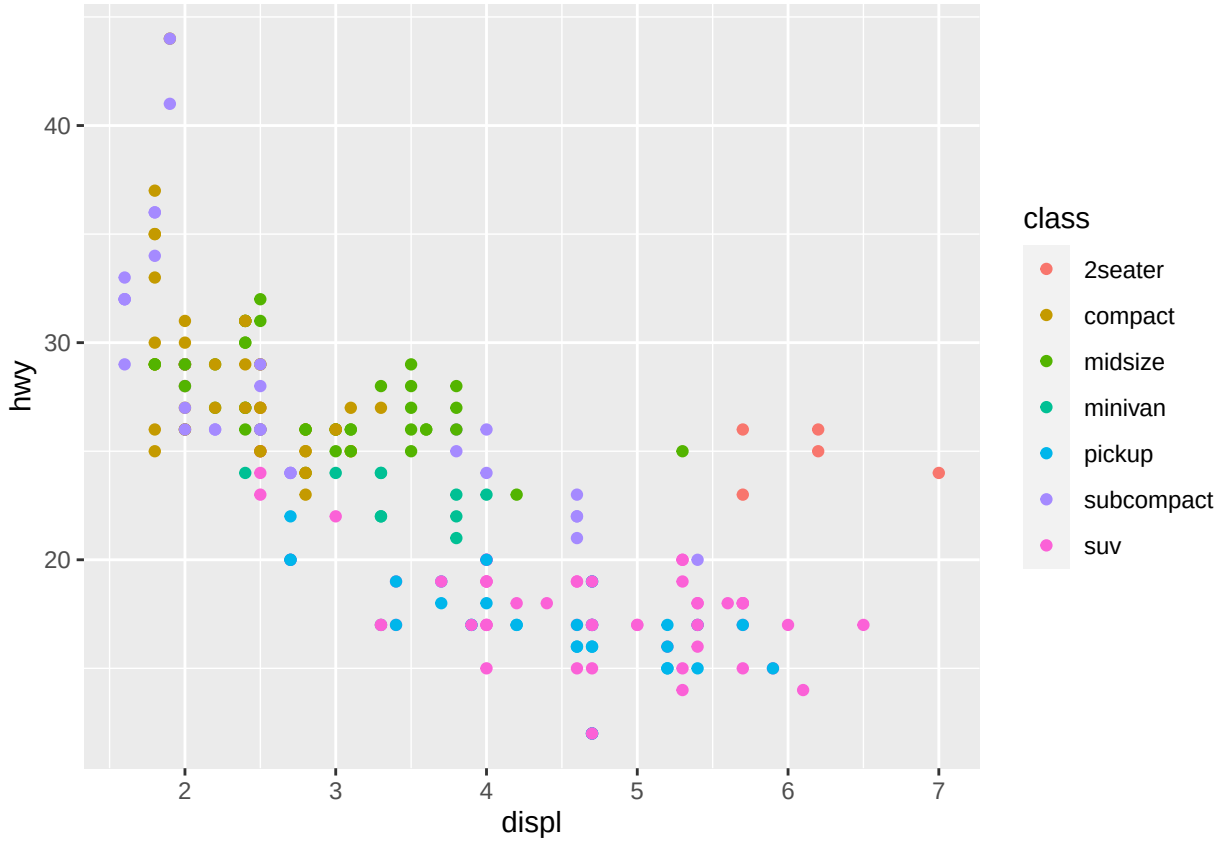

Adding a size aesthetic:

```
mpg %>%                          # pipe in the data
  ggplot() +                     # create an empty ggplot
  geom_point(mapping =           # add scatter plot
               aes(x = displ,    # x axis location of points
                   y = hwy,      # y axis location of points
                   size = class))  # size of points
```

```
## Warning: Using size for a discrete variable is not advised.
```

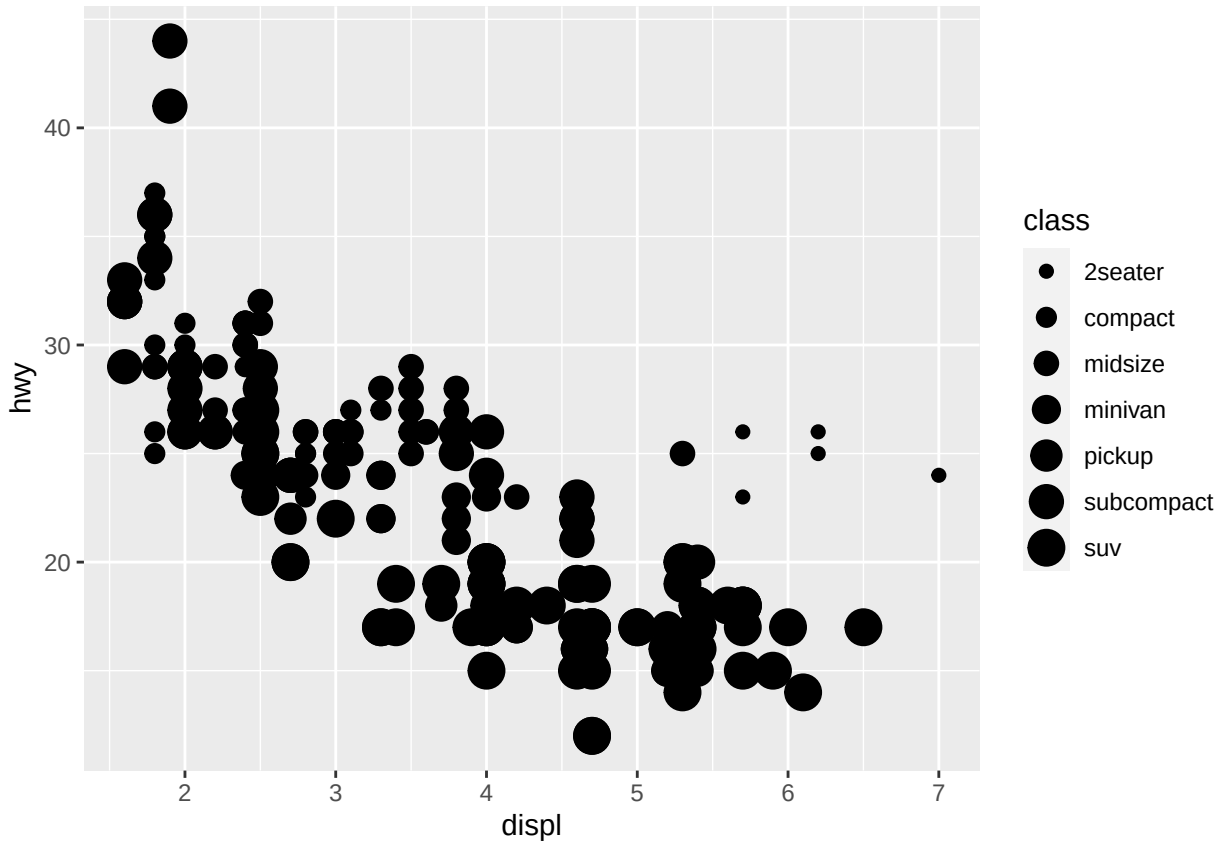

Adding a shape aesthetic:

```
mpg %>%                           # pipe in the data
  ggplot() +                      # create an empty ggplot
  geom_point(mapping =            # add scatter plot
               aes(x = displ,     # x axis location of points
                   y = hwy,       # y axis location of points
                   shape = class)) # size of points
```

```
## Warning: The shape palette can deal with a maximum of 6 discrete values because
## more than 6 becomes difficult to discriminate; you have 7. Consider
## specifying shapes manually if you must have them.
```

```
## Warning: Removed 62 rows containing missing values (geom_point).
```

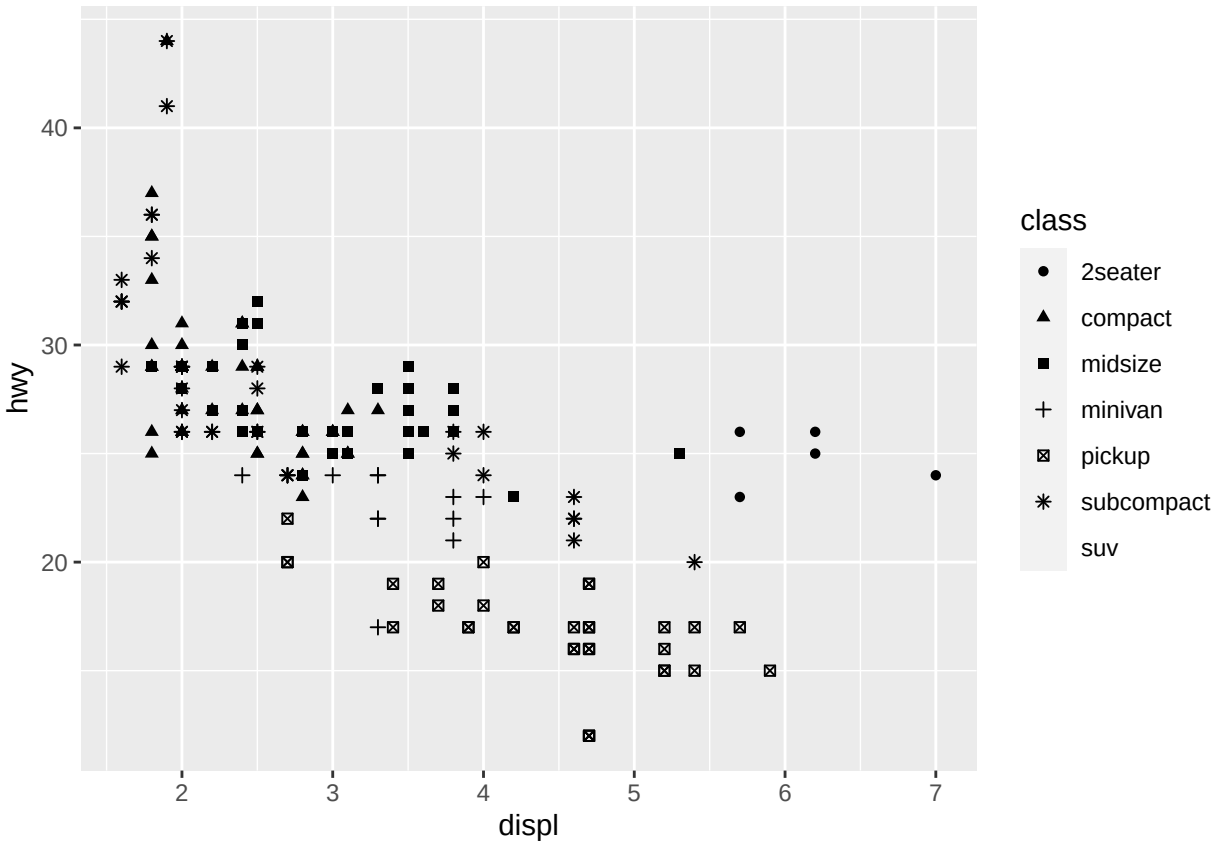

Adding a transparency aesthetic:

```
mpg %>%                           # pipe in the data
  ggplot() +                      # create an empty ggplot
  geom_point(mapping =            # add scatter plot
               aes(x = displ,     # x axis location of points
                   y = hwy,       # y axis location of points
                   alpha = class)) # transparency of points
```

```
## Warning: Using alpha for a discrete variable is not advised.
```

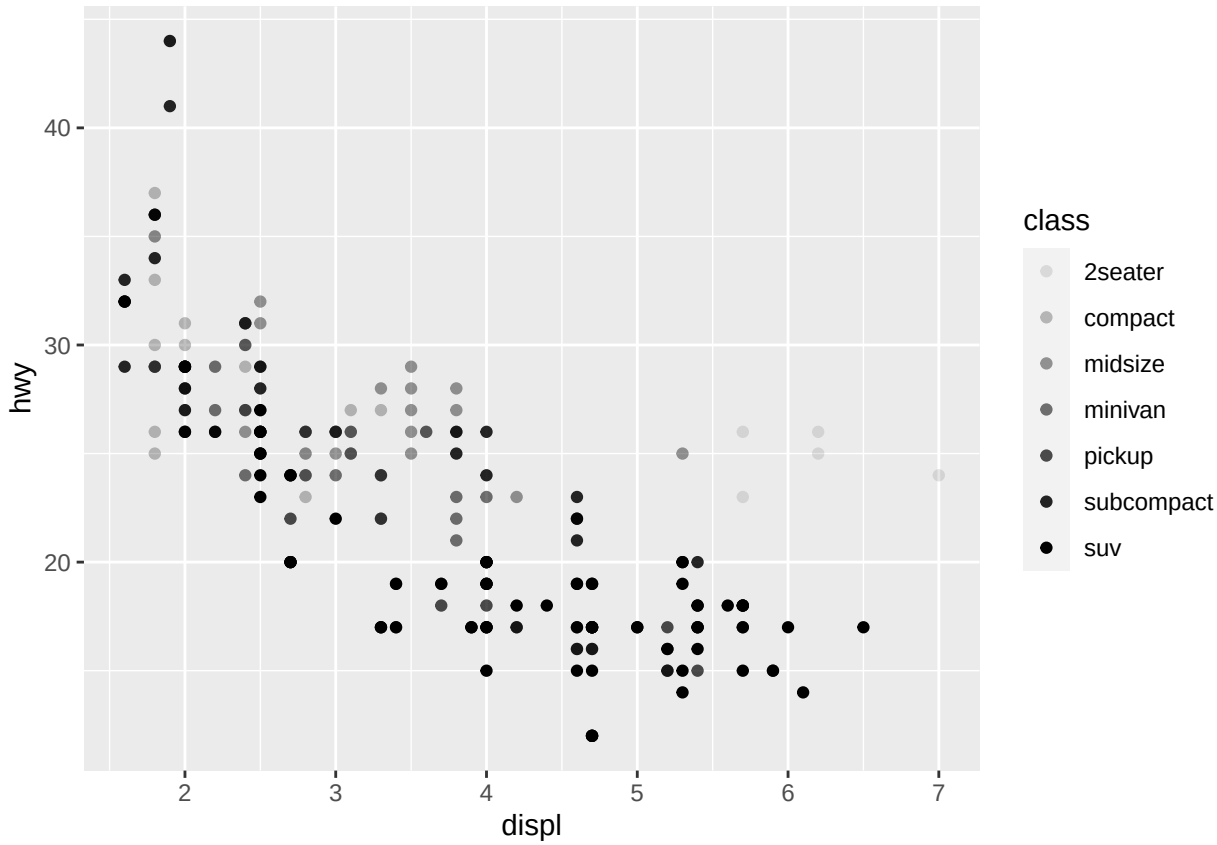

Specifying an aesthetic manually, instead of mapping from a variable:

```
mpg %>%                             # pipe in the data
  ggplot() +                        # create an empty ggplot
  geom_point(mapping =              # add scatter plot
               aes(x = displ,       # x axis location of points
                   y = hwy),        # y axis location of points
             color = "blue")        # color of points (outside of aes)
```

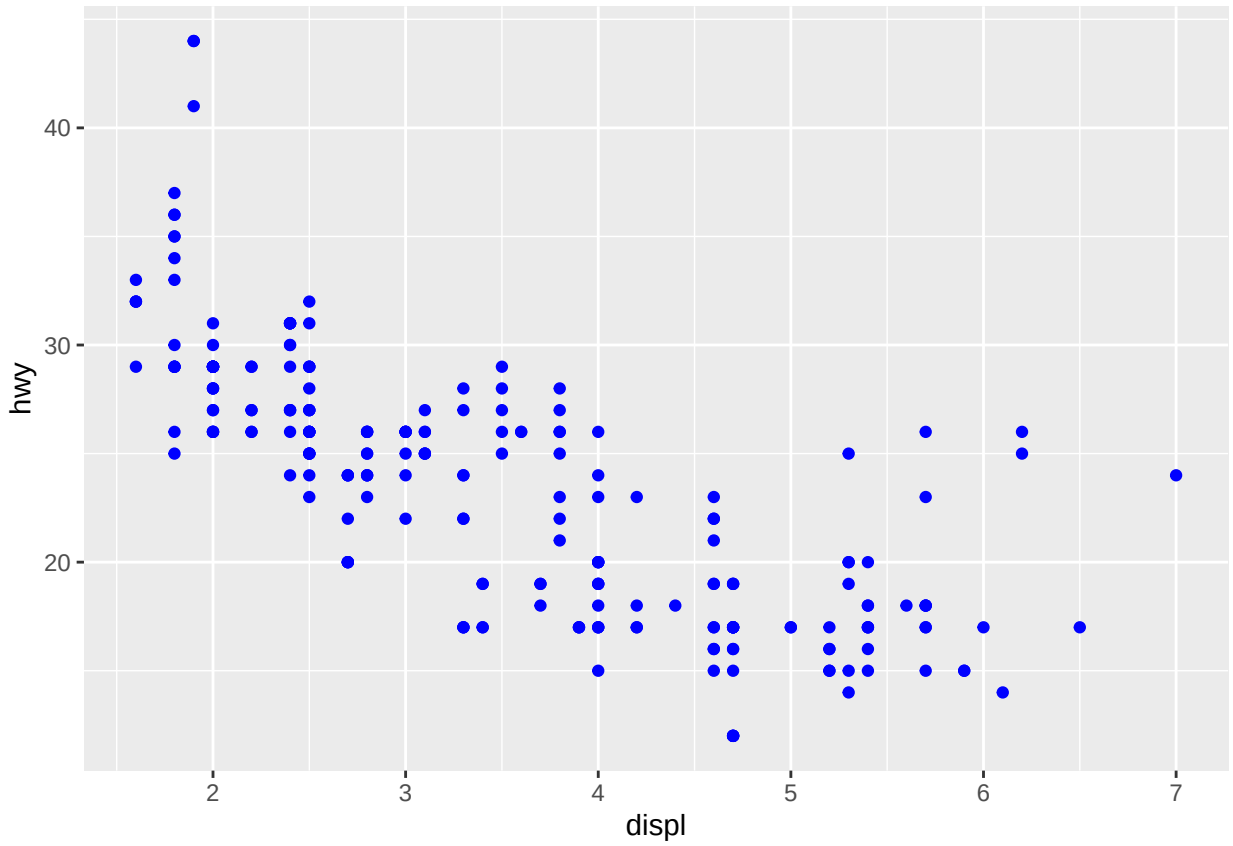

1.3 Exercises

1. What's gone wrong with this code? Why are the points not blue?

```
mpg %>%
  ggplot() +
  geom_point(mapping = aes(x = displ, y = hwy, color = "blue"))
```

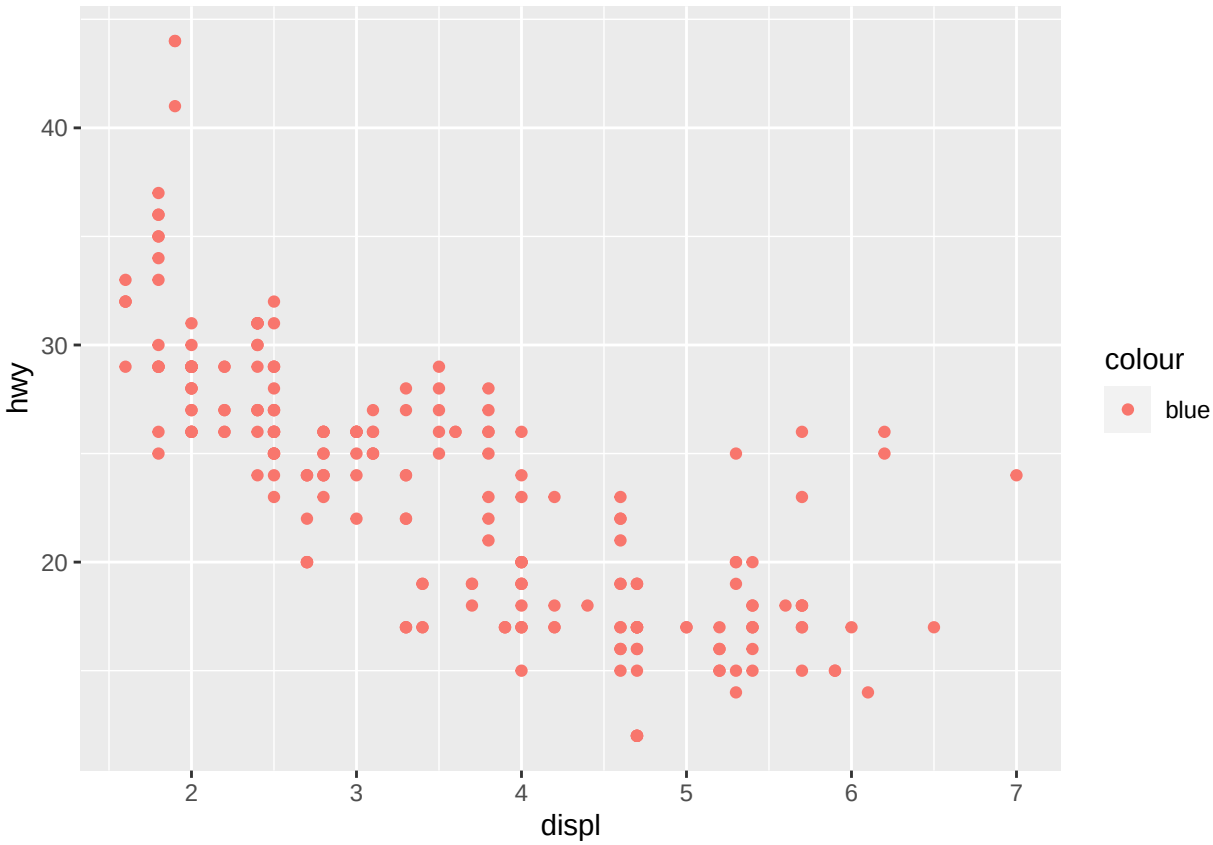

2. Map a continuous variable to `color`, `size`, and `shape`. How do these aesthetics behave differently for categorical vs. continuous variables?

3. What happens if you map the same variable to multiple aesthetics?

4. What happens if you map an aesthetic to something other than a variable name, like `aes(colour = displ < 5)`? Note, you'll also need to specify `x` and `y`.

1.4 Facets

To facet your plot by a single variable, use `facet_wrap()`. The first argument of `facet_wrap()` should be a formula, which you create with `~` followed by a variable name (here "formula" is the name of a data structure in R, not a synonym for "equation"). The variable that you pass to `facet_wrap()` should be discrete.

```
mpg %>%                          # pipe in the data
  ggplot() +                     # create empty ggplot
  geom_point(mapping =           # create scatter plot
               aes(x = displ,    # map displ to x axis location
                   y = hwy)) +   # map hwy to y axis location
  facet_wrap(~ class,            # split into facets based on class
             nrow = 2)           # have two rows of facets
```

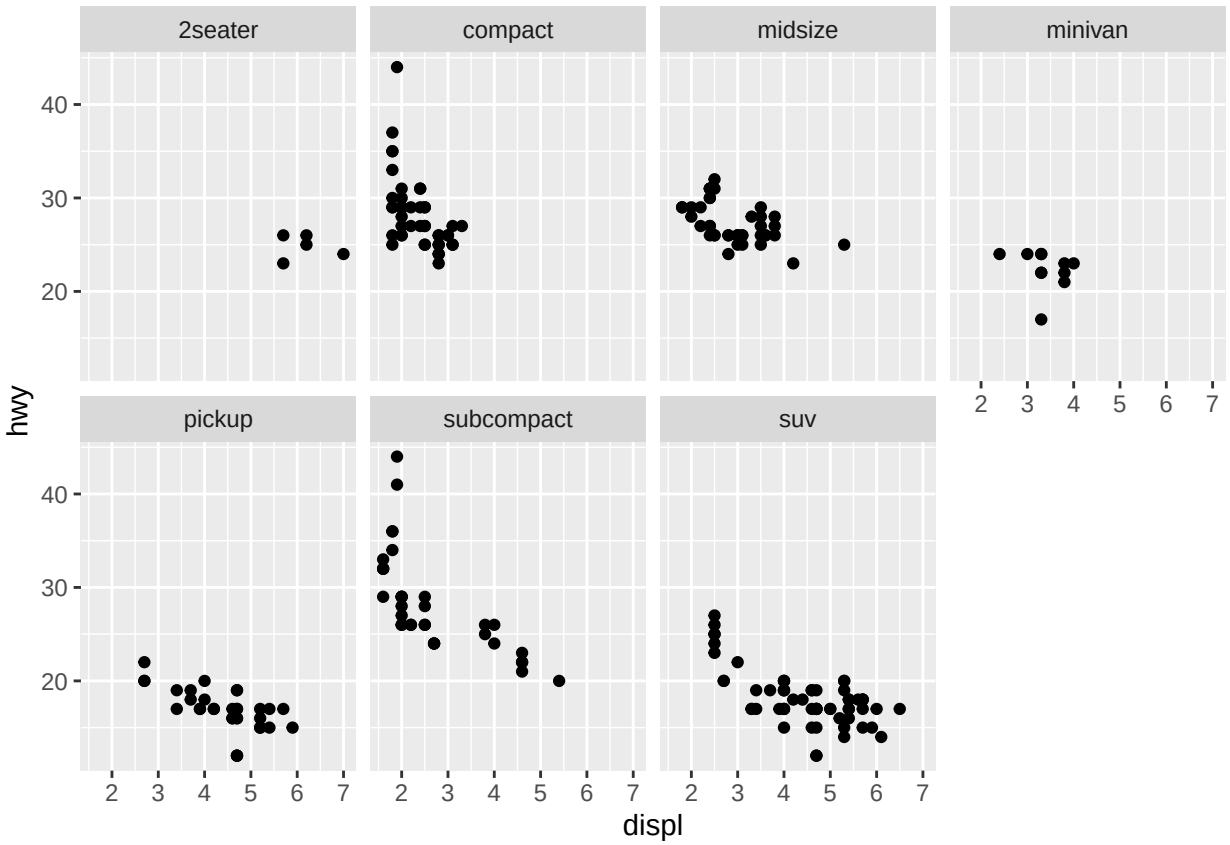

To facet your plot on the combination of two variables, add `facet_grid()` to your plot call. The first argument of `facet_grid()` is also a formula. This time the formula should contain two variable names separated by a `~`.

```
mpg %>%
  ggplot() +
  geom_point(mapping =
               aes(x = displ, y = hwy)) +
  facet_grid(drv ~ cyl)                      # facet on drv and cyl
```

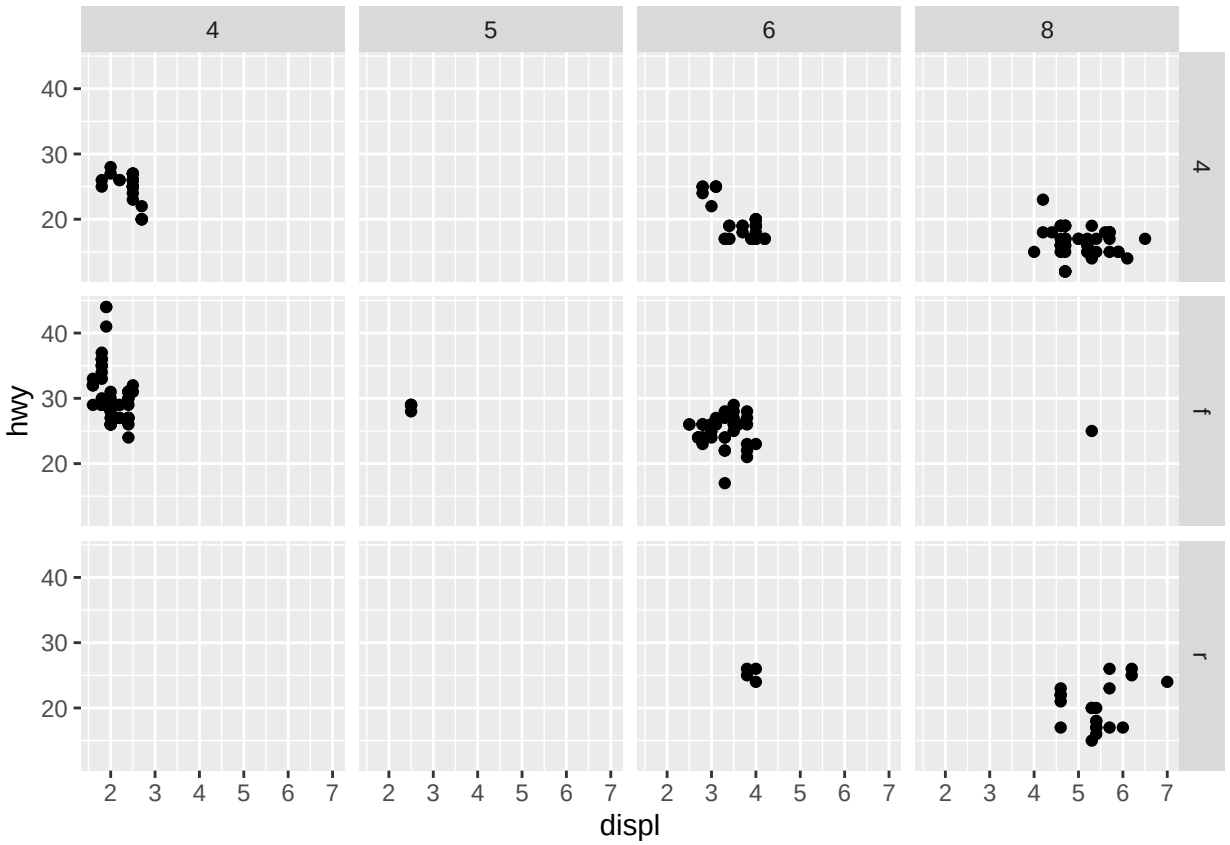

If you prefer to not facet in the rows or columns dimension, use a `.` instead of a variable name, e.g. `+ facet_grid(. ~ cyl)`.

1.5 Exercises

1. What happens if you facet on a continuous variable?
2. Why are there empty facets in the plot with `facet_grid(drv ~ cyl)`?
3. Read `?facet_wrap`. What does `nrow` do? What does `ncol` do? Why doesn't `facet_grid()` have `nrow` and `ncol` arguments?

1.6 geoms

To visualize the distribution of a quantitative variable:

```
mpg %>%
  ggplot() +
  geom_histogram(aes(x = hwy)) # we usually drop "mapping ="
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

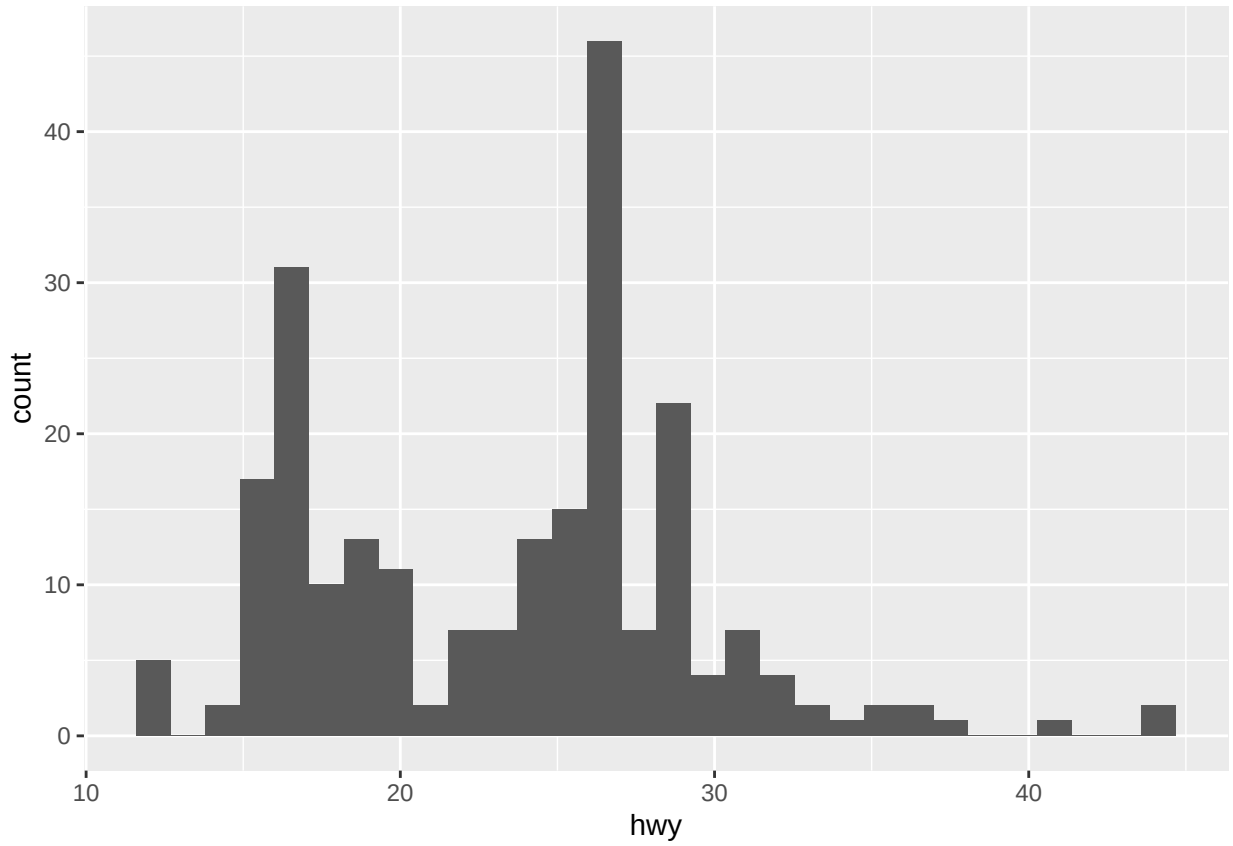

To visualize the distribution of a categorical variable:

```
# bar plot
mpg %>%
  ggplot() +
  geom_bar(aes(x = class))
```

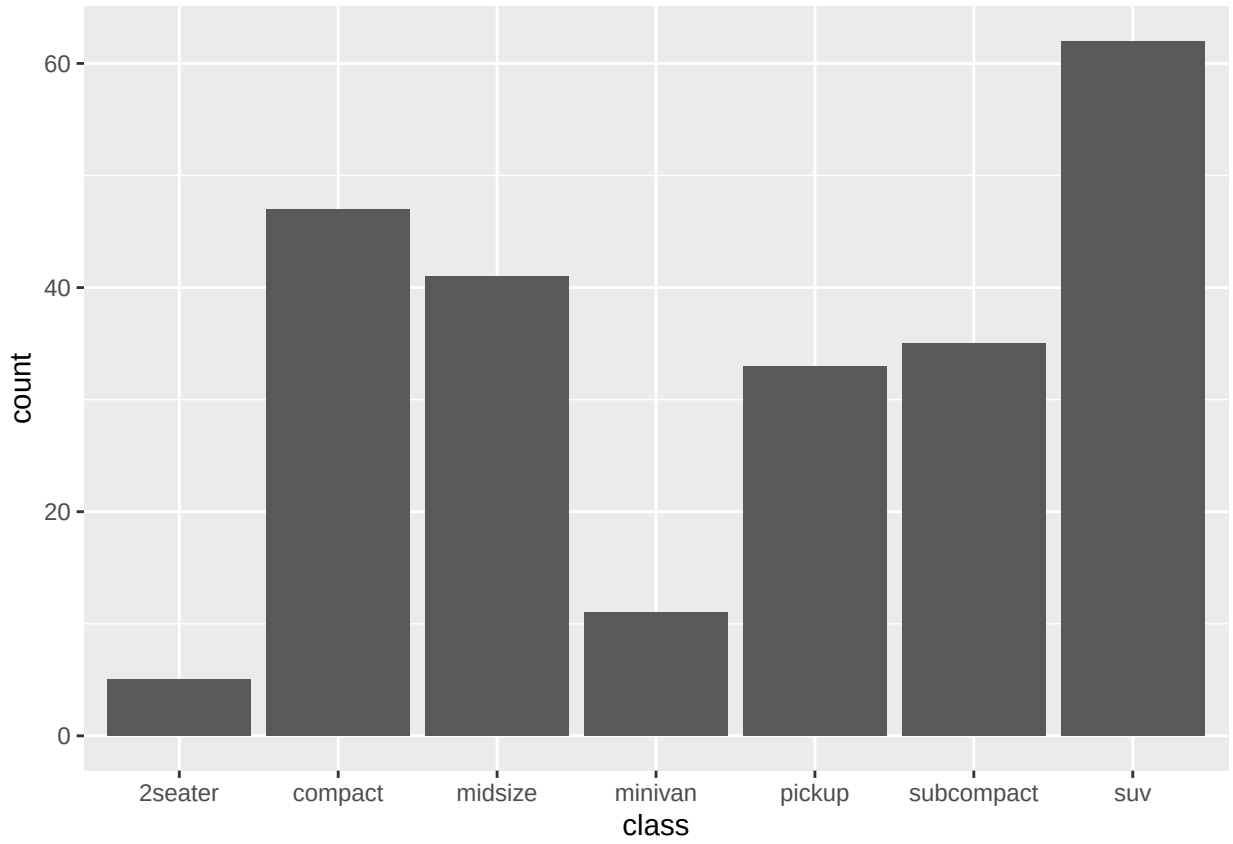


```
# better to reorder class variable
mpg %>%
  ggplot() +
  geom_bar(aes(x = fct_infreq(class)))
```

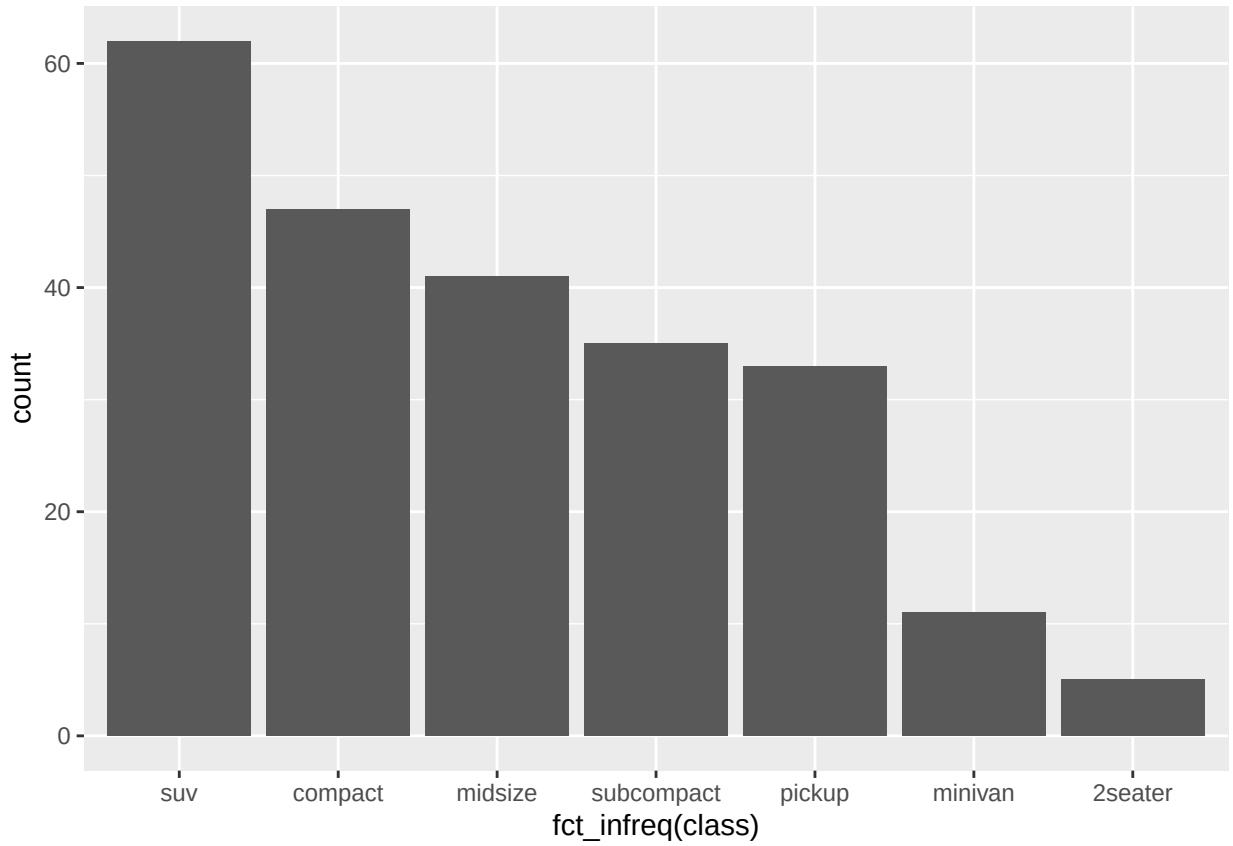

To visualize the relationship between a quantitative and categorical variable:

```
# boxplot
mpg %>%
  ggplot() +
  geom_boxplot(aes(x = class, y = hwy))
```

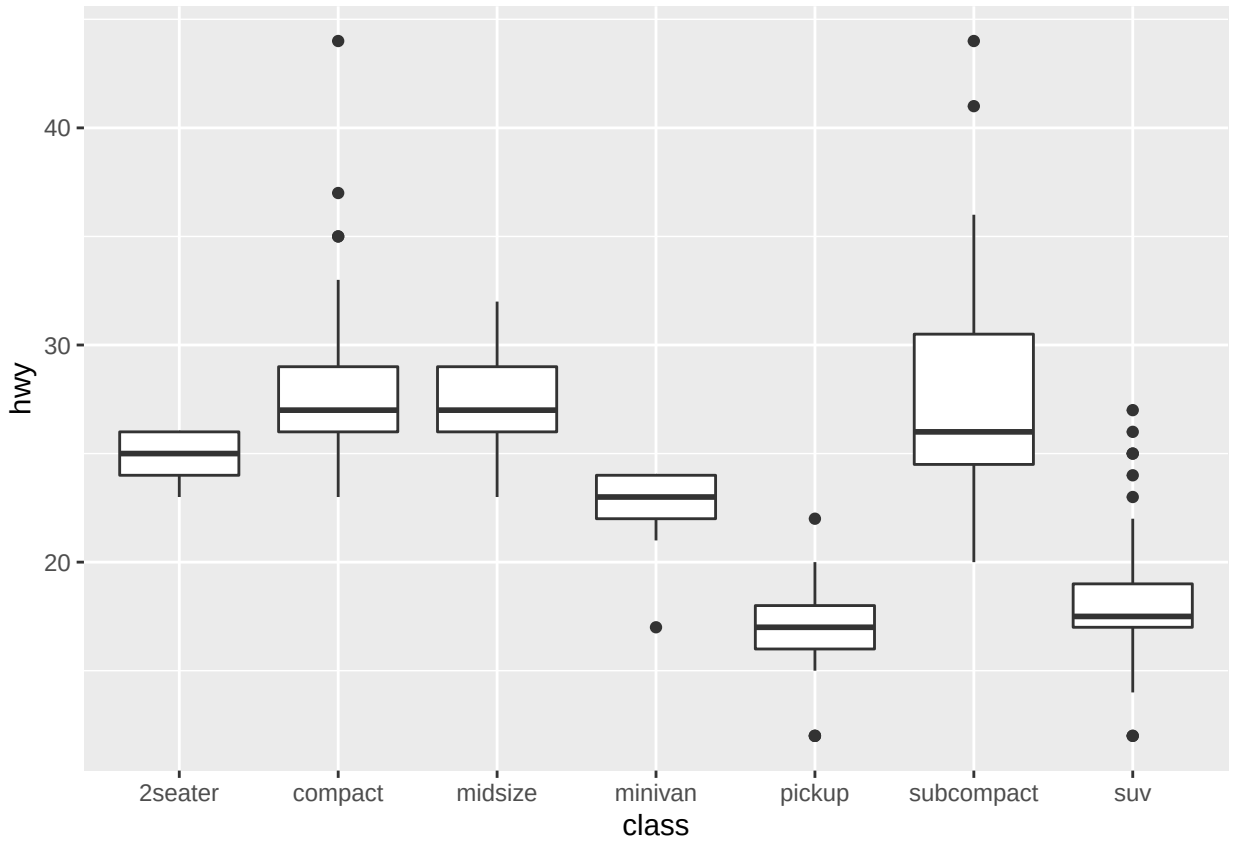


```
# better to reorder factors based on hwy
mpg %>%
  ggplot() +
  geom_boxplot(aes(x = fct_reorder(class, hwy), y = hwy))
```

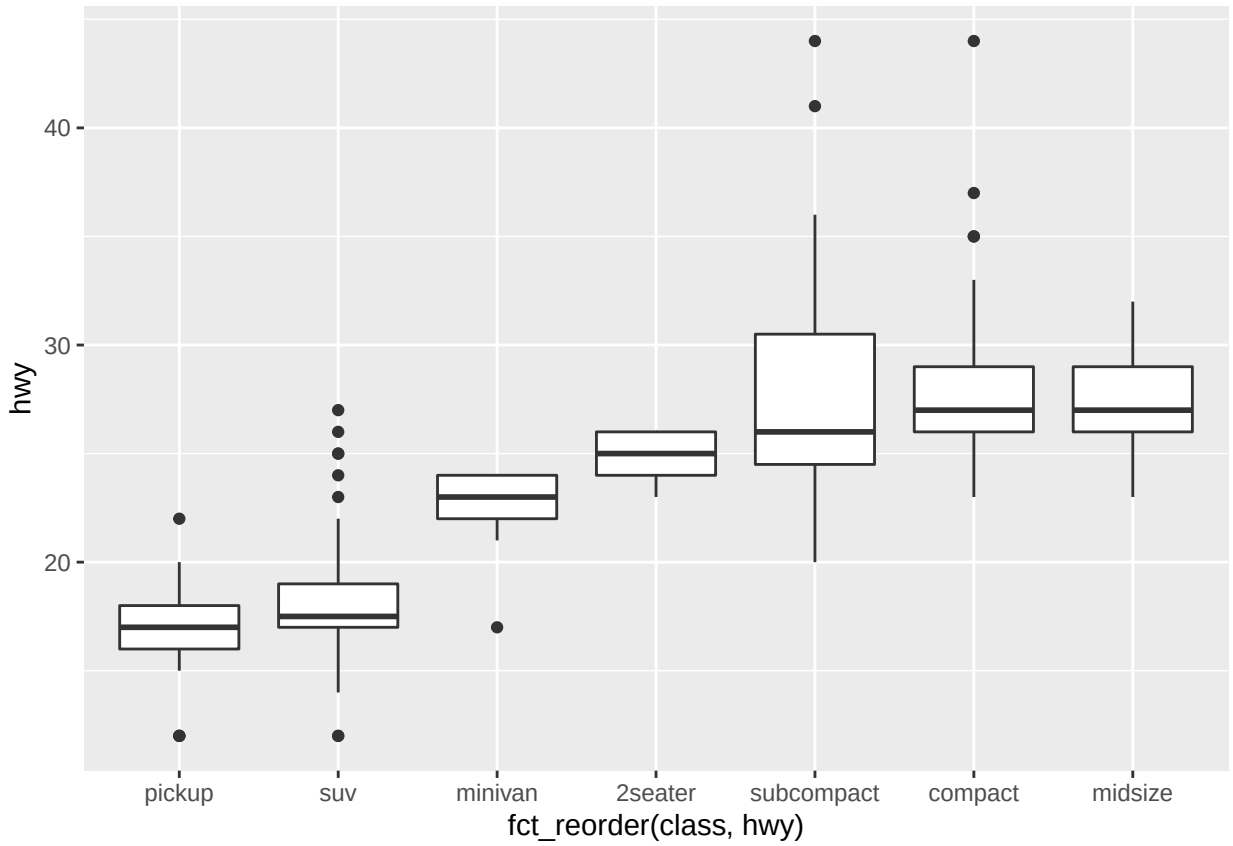

To add a smooth curve:

```
mpg %>%
  ggplot() +
  geom_point(aes(x = displ, y = hwy)) +              # create scatter plot
  geom_smooth(aes(x = displ, y = hwy), se = FALSE)   # add smooth curve
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

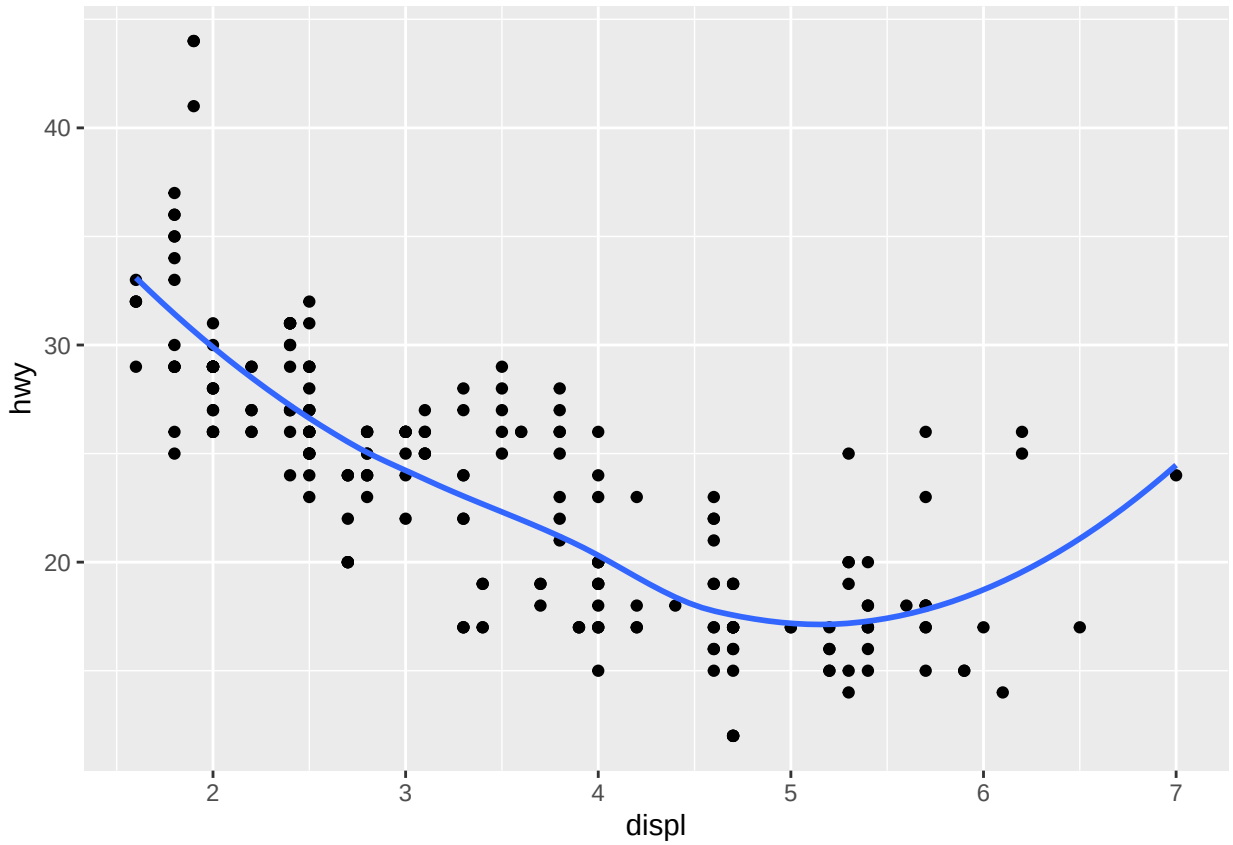

Can instead set aesthetic mapping *globally*:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy)) + # set aesthetic mapping globally
  geom_point() +                    # add scatter plot
  geom_smooth(se = FALSE)           # add smooth curve
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

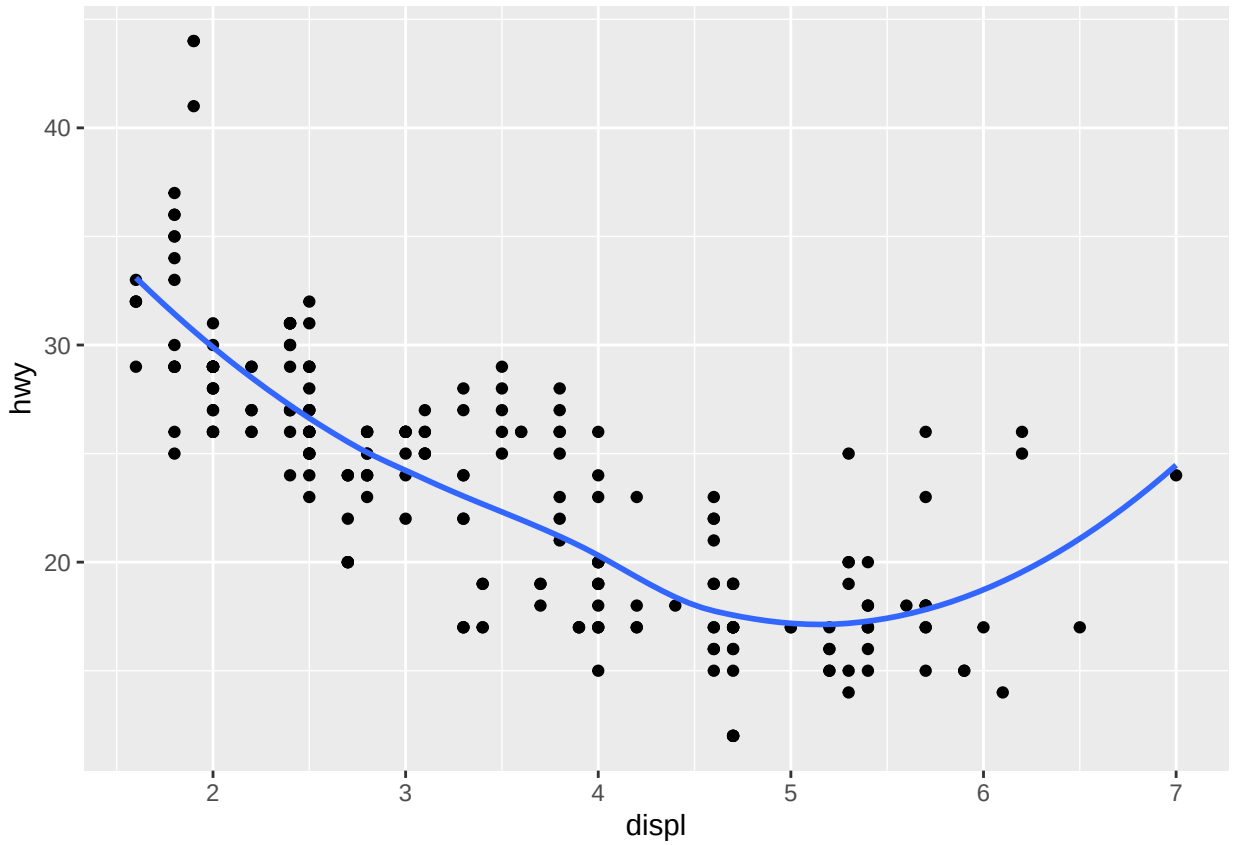

Can set some aesthetic mappings globally and some locally:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy)) + # set x and y aesthetic mappings globally
  geom_point(aes(color = class)) +  # set color aesthetic mapping locally
  geom_smooth(se = FALSE)           # add smooth curve
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

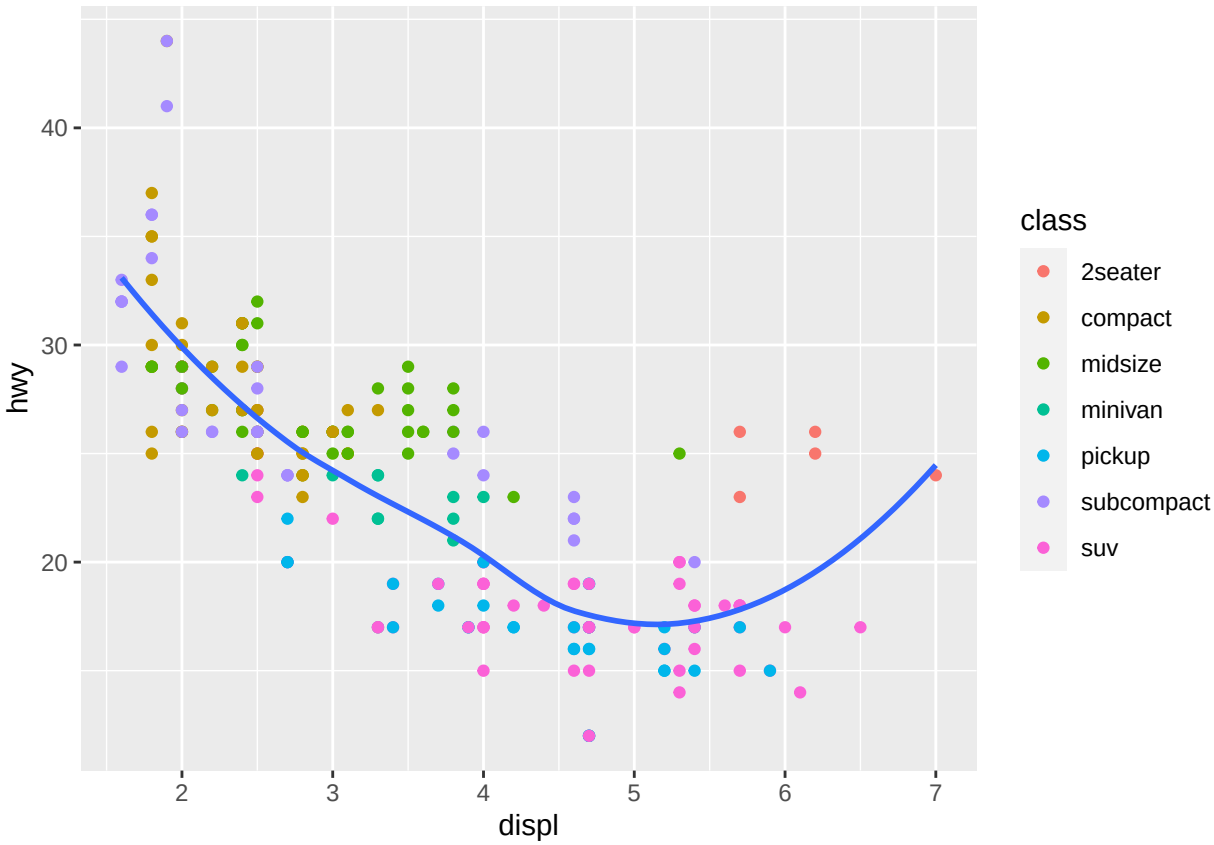

If we had set `color` aesthetic globally, ggplot would separate smooth curves by `class`:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy, color = class)) + # set x, y, color aesthetic
  geom_point() +                                   #  mappings globally
  geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : span too small. fewer data values than degrees of freedom.

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at 5.6935

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 0.5065

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 0

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : There are other near singularities as well. 0.65044

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at 4.008

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 0.708

## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 0
```

```
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : There are other near singularities as well. 0.25
```

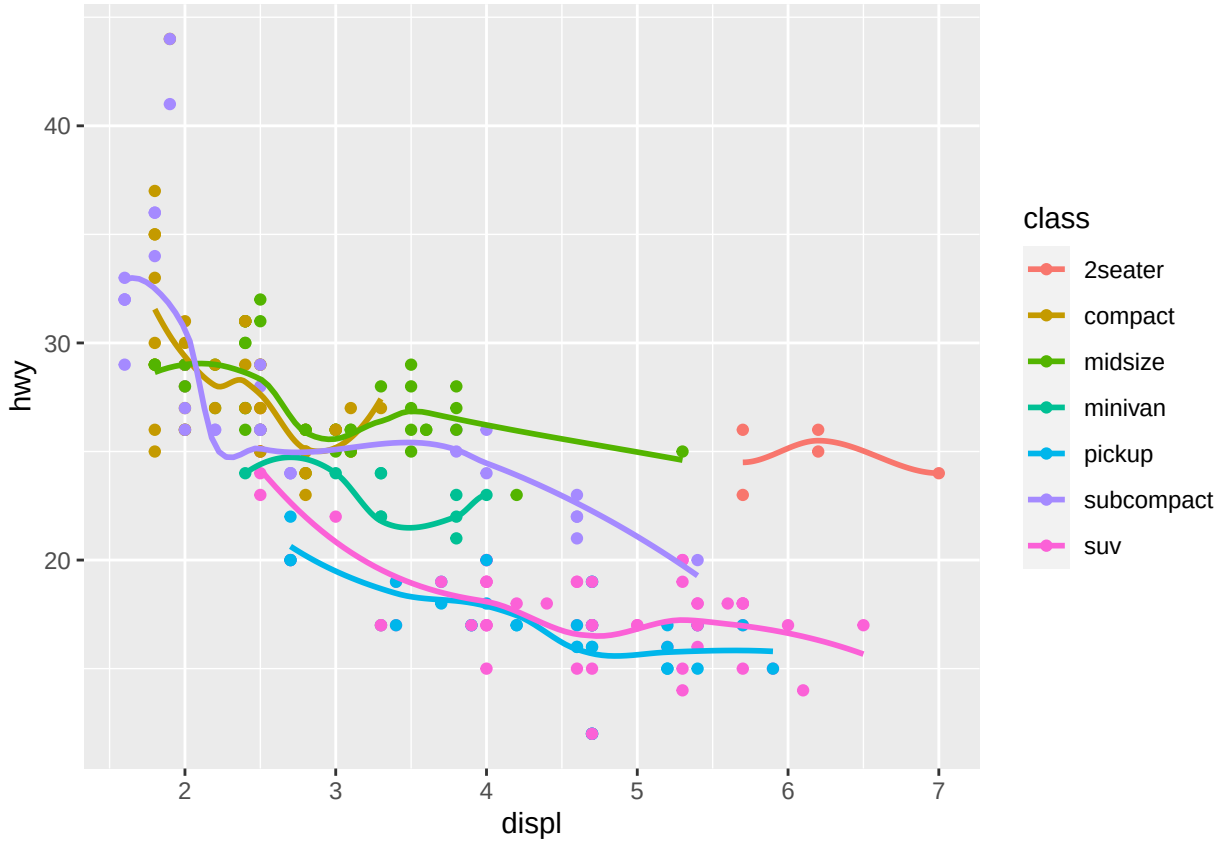

Add horizontal, vertical, and/or oblique lines:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy)) + # set aesthetic mapping globally
  geom_point() +                    # add scatter plot
  geom_smooth(se = FALSE) +         # add smooth curve
  geom_vline(xintercept = 4,        # add vertical line with x-intercept 4
             linetype = "dashed",
             color = "red") +
  geom_hline(yintercept = 20,       # add horizontal line with y-intercept 20
             linetype = "dashed",
             color = "red") +
  geom_abline(slope = 10,           # add oblique line with slope 10
              intercept = 2,        #  and y-intercept 2
              linetype = "solid",
              color = "red")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

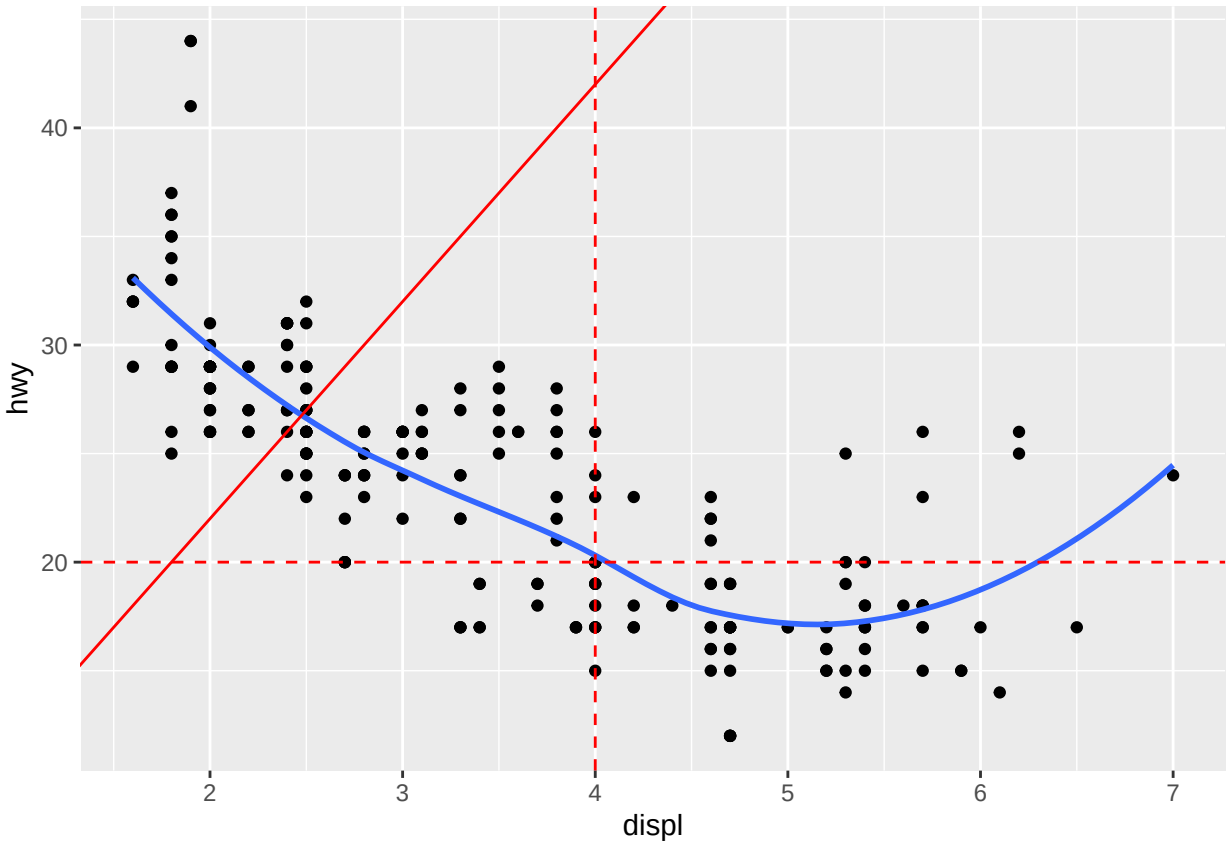

1.7 Themes

Different pre-set "themes" give plots different appearances. A theme I prefer to use is `theme_bw()`:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy)) + # set aesthetic mapping globally
  geom_point() +                    # add scatter plot
  geom_smooth(se = FALSE) +         # add smooth curve
  geom_vline(xintercept = 4,        # add vertical line with x-intercept 4
             linetype = "dashed",
             color = "red") +
  geom_hline(yintercept = 20,       # add horizontal line with y-intercept 20
             linetype = "dashed",
             color = "red") +
  geom_abline(slope = 10,           # add oblique line with slope 10
              intercept = 2,        #  and y-intercept 2
              linetype = "solid",
              color = "red") +
  theme_bw()                        # add classy bw theme
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

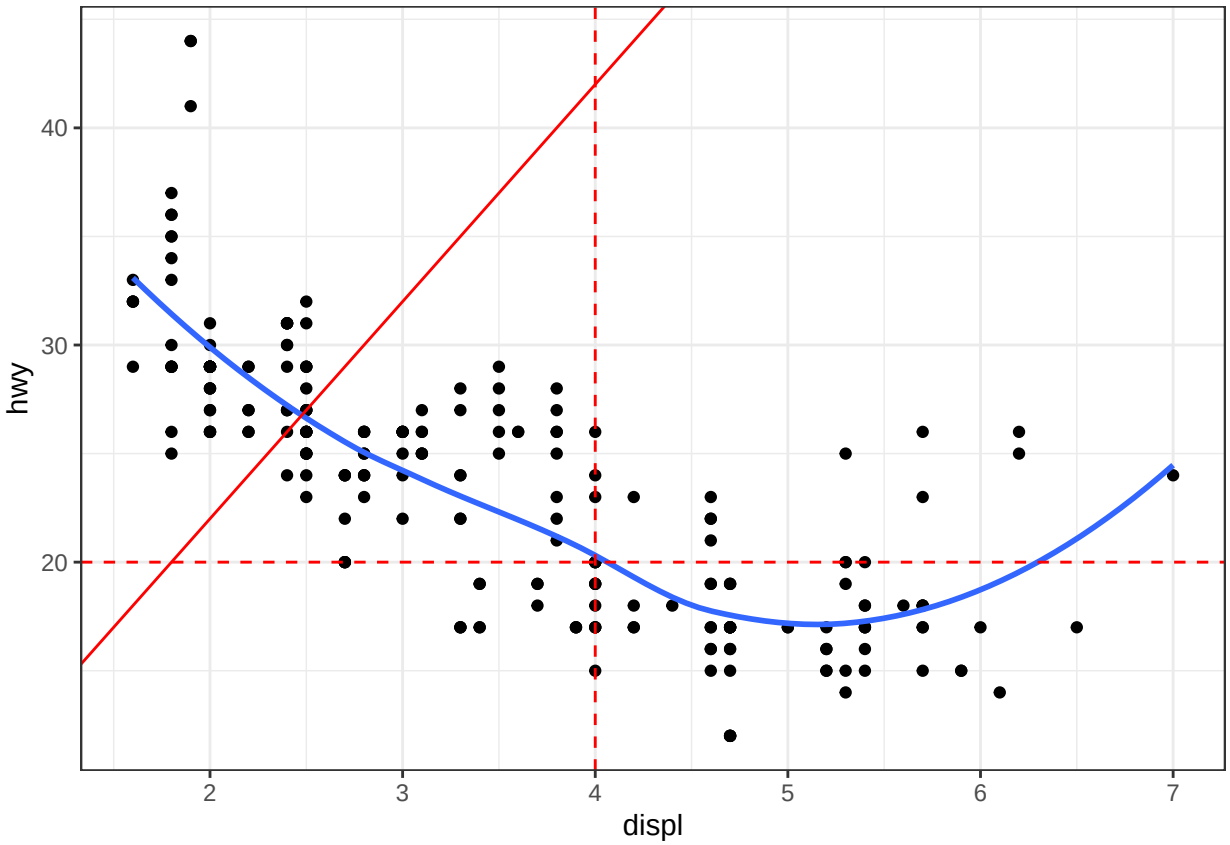

1.8 High-quality figures for communication

See preparing-reports.pdf. Part of your homework will be graded on presentation quality.

2 Data transformation

The `dplyr` package (another core member of the `tidyverse`) facilitates manipulation of data. This includes key operations:

- Pick observations by their values (`filter()`).
- Reorder the rows (`arrange()`).
- Pick variables by their names (`select()`).
- Create new variables with functions of existing variables (`mutate()`).
- Collapse many values down to a single summary (`summarise()`).

These can all be used in conjunction with `group_by()` which changes the scope of each function from operating on the entire dataset to operating on it group-by-group. These six functions provide the verbs for a language of data manipulation.

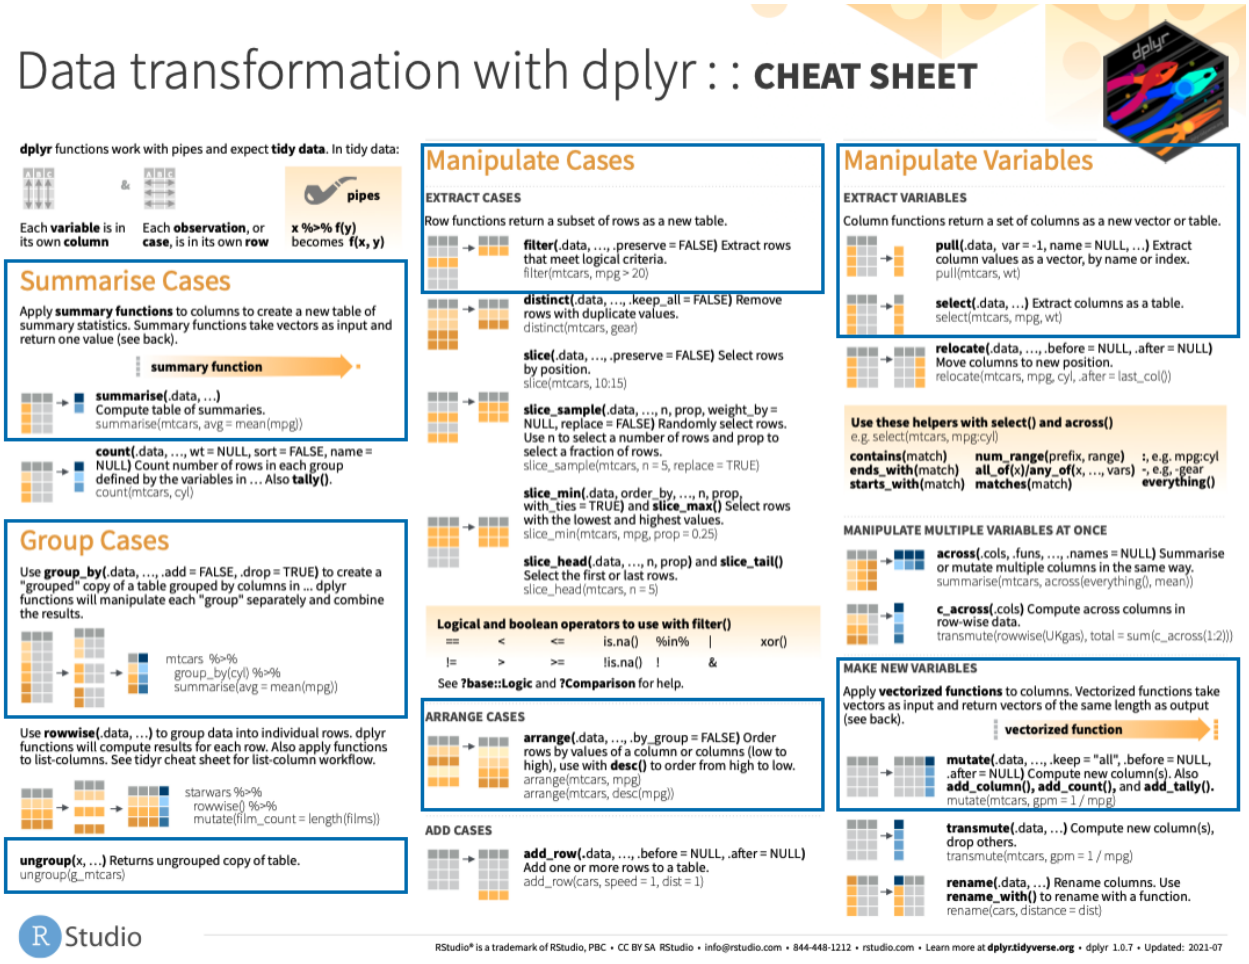

Figure 4: Image source: https://www.rstudio.com/resources/cheatsheets/

2.1 Filter rows with `filter()`

```
mpg %>%
  filter(class == "compact")
```

```
## # A tibble: 47 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 audi         a4           1.8  1999     4 auto~ f        18    29 p     comp~
##  2 audi         a4           1.8  1999     4 manu~ f        21    29 p     comp~
##  3 audi         a4           2    2008     4 manu~ f        20    31 p     comp~
##  4 audi         a4           2    2008     4 auto~ f        21    30 p     comp~
##  5 audi         a4           2.8  1999     6 auto~ f        16    26 p     comp~
##  6 audi         a4           2.8  1999     6 manu~ f        18    26 p     comp~
##  7 audi         a4           3.1  2008     6 auto~ f        18    27 p     comp~
##  8 audi         a4 quattro   1.8  1999     4 manu~ 4        18    26 p     comp~
##  9 audi         a4 quattro   1.8  1999     4 auto~ 4        16    25 p     comp~
## 10 audi         a4 quattro   2    2008     4 manu~ 4        20    28 p     comp~
## # ... with 37 more rows
```

```
mpg %>%
  filter(class %in% c("compact", "2seater"))
```

```
## # A tibble: 52 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 audi         a4           1.8  1999     4 auto~ f        18    29 p     comp~
##  2 audi         a4           1.8  1999     4 manu~ f        21    29 p     comp~
##  3 audi         a4           2    2008     4 manu~ f        20    31 p     comp~
##  4 audi         a4           2    2008     4 auto~ f        21    30 p     comp~
##  5 audi         a4           2.8  1999     6 auto~ f        16    26 p     comp~
##  6 audi         a4           2.8  1999     6 manu~ f        18    26 p     comp~
##  7 audi         a4           3.1  2008     6 auto~ f        18    27 p     comp~
##  8 audi         a4 quattro   1.8  1999     4 manu~ 4        18    26 p     comp~
##  9 audi         a4 quattro   1.8  1999     4 auto~ 4        16    25 p     comp~
## 10 audi         a4 quattro   2    2008     4 manu~ 4        20    28 p     comp~
## # ... with 42 more rows
```

```
mpg %>%
  filter(class %in% c("compact", "2seater") & year > 2000) # & means "and"
```

```
## # A tibble: 25 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 audi         a4           2    2008     4 manu~ f        20    31 p     comp~
##  2 audi         a4           2    2008     4 auto~ f        21    30 p     comp~
##  3 audi         a4           3.1  2008     6 auto~ f        18    27 p     comp~
##  4 audi         a4 quattro   2    2008     4 manu~ 4        20    28 p     comp~
##  5 audi         a4 quattro   2    2008     4 auto~ 4        19    27 p     comp~
##  6 audi         a4 quattro   3.1  2008     6 auto~ 4        17    25 p     comp~
##  7 audi         a4 quattro   3.1  2008     6 manu~ 4        15    25 p     comp~
##  8 chevrolet    corvette     6.2  2008     8 manu~ r        16    26 p     2sea~
##  9 chevrolet    corvette     6.2  2008     8 auto~ r        15    25 p     2sea~
## 10 chevrolet    corvette     7    2008     8 manu~ r        15    24 p     2sea~
## # ... with 15 more rows
```

2.2 Exercises

1. Find all cars manufactured by a Japanese company (`Honda`, `Toyota`, `Nissan`, `Subaru`). How many such cars are there in these data?
2. Find all cars whose highway fuel efficiency (`hwy`) exceeded their city fuel efficiency (`cty`) by at least a factor of 1.5. How many such cars are there in these data?

```
mpg %>%
  filter(class %in% c("compact", "2seater") | year > 2000) # | means "or"
```

```
## # A tibble: 144 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 audi         a4           1.8  1999     4 auto~ f        18    29 p     comp~
##  2 audi         a4           1.8  1999     4 manu~ f        21    29 p     comp~
##  3 audi         a4           2    2008     4 manu~ f        20    31 p     comp~
##  4 audi         a4           2    2008     4 auto~ f        21    30 p     comp~
##  5 audi         a4           2.8  1999     6 auto~ f        16    26 p     comp~
##  6 audi         a4           2.8  1999     6 manu~ f        18    26 p     comp~
##  7 audi         a4           3.1  2008     6 auto~ f        18    27 p     comp~
##  8 audi         a4 quattro   1.8  1999     4 manu~ 4        18    26 p     comp~
##  9 audi         a4 quattro   1.8  1999     4 auto~ 4        16    25 p     comp~
## 10 audi         a4 quattro   2    2008     4 manu~ 4        20    28 p     comp~
## # ... with 134 more rows
```

2.3 Arrange rows with `arrange()`

You can sort the rows of a tibble according to the values of a certain variable:

```
mpg %>% arrange(hwy)
```

```
## # A tibble: 234 x 11
##    manufacturer model      displ  year   cyl trans  drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr>  <chr> <int> <int> <chr> <chr>
##  1 dodge        dakota p~   4.7  2008     8 auto(~ 4         9    12 e     pick~
##  2 dodge        durango ~   4.7  2008     8 auto(~ 4         9    12 e     suv
##  3 dodge        ram 1500~   4.7  2008     8 auto(~ 4         9    12 e     pick~
##  4 dodge        ram 1500~   4.7  2008     8 manua~ 4         9    12 e     pick~
##  5 jeep         grand ch~   4.7  2008     8 auto(~ 4         9    12 e     suv
##  6 chevrolet    k1500 ta~   5.3  2008     8 auto(~ 4        11    14 e     suv
##  7 jeep         grand ch~   6.1  2008     8 auto(~ 4        11    14 p     suv
##  8 chevrolet    c1500 su~   5.3  2008     8 auto(~ r        11    15 e     suv
##  9 chevrolet    k1500 ta~   5.7  1999     8 auto(~ 4        11    15 r     suv
## 10 dodge        dakota p~   5.2  1999     8 auto(~ 4        11    15 r     pick~
## # ... with 224 more rows
```

Or in descending order:

```
mpg %>% arrange(desc(hwy))
```

```
## # A tibble: 234 x 11
##    manufacturer model      displ  year   cyl trans drv     cty   hwy fl    class
##    <chr>        <chr>      <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
##  1 volkswagen   jetta        1.9  1999     4 manu~ f        33    44 d     comp~
##  2 volkswagen   new beetle   1.9  1999     4 manu~ f        35    44 d     subc~
##  3 volkswagen   new beetle   1.9  1999     4 auto~ f        29    41 d     subc~
##  4 toyota       corolla      1.8  2008     4 manu~ f        28    37 r     comp~
```

```
##  5 honda        civic        1.8  2008     4 auto~ f        25    36 r     subc~
##  6 honda        civic        1.8  2008     4 auto~ f        24    36 c     subc~
##  7 toyota       corolla      1.8  1999     4 manu~ f        26    35 r     comp~
##  8 toyota       corolla      1.8  2008     4 auto~ f        26    35 r     comp~
##  9 honda        civic        1.8  2008     4 manu~ f        26    34 r     subc~
## 10 honda        civic        1.6  1999     4 manu~ f        28    33 r     subc~
## # ... with 224 more rows
```

Which car had the best highway fuel efficiency?

2.4 Select columns with `select()`

Select columns:

```
mpg %>% select(manufacturer, model, year)
```

```
## # A tibble: 234 x 3
##    manufacturer model       year
##    <chr>        <chr>      <int>
##  1 audi         a4          1999
##  2 audi         a4          1999
##  3 audi         a4          2008
##  4 audi         a4          2008
##  5 audi         a4          1999
##  6 audi         a4          1999
##  7 audi         a4          2008
##  8 audi         a4 quattro  1999
##  9 audi         a4 quattro  1999
## 10 audi         a4 quattro  2008
## # ... with 224 more rows
```

De-select columns:

```
mpg %>% select(-manufacturer, -model, -year)
```

```
## # A tibble: 234 x 8
##    displ   cyl trans      drv     cty   hwy fl    class
##    <dbl> <int> <chr>      <chr> <int> <int> <chr> <chr>
##  1   1.8     4 auto(l5)   f        18    29 p     compact
##  2   1.8     4 manual(m5) f        21    29 p     compact
##  3   2       4 manual(m6) f        20    31 p     compact
##  4   2       4 auto(av)   f        21    30 p     compact
##  5   2.8     6 auto(l5)   f        16    26 p     compact
##  6   2.8     6 manual(m5) f        18    26 p     compact
##  7   3.1     6 auto(av)   f        18    27 p     compact
##  8   1.8     4 manual(m5) 4        18    26 p     compact
##  9   1.8     4 auto(l5)   4        16    25 p     compact
## 10   2       4 manual(m6) 4        20    28 p     compact
## # ... with 224 more rows
```

2.5 Add new variables with `mutate()`

```
mpg_small = mpg %>% select(manufacturer, model, year, cty, hwy)
mpg_small %>%
  mutate(hwy_boost = hwy - cty,
         japanese = manufacturer %in% c("Honda", "Toyota", "Nissan", "Subaru"))
```

```
## # A tibble: 234 x 7
##    manufacturer model       year   cty   hwy hwy_boost japanese
##    <chr>        <chr>      <int> <int> <int>     <int> <lgl>
##  1 audi         a4          1999    18    29        11 FALSE
##  2 audi         a4          1999    21    29         8 FALSE
##  3 audi         a4          2008    20    31        11 FALSE
##  4 audi         a4          2008    21    30         9 FALSE
##  5 audi         a4          1999    16    26        10 FALSE
##  6 audi         a4          1999    18    26         8 FALSE
##  7 audi         a4          2008    18    27         9 FALSE
##  8 audi         a4 quattro  1999    18    26         8 FALSE
##  9 audi         a4 quattro  1999    16    25         9 FALSE
## 10 audi         a4 quattro  2008    20    28         8 FALSE
## # ... with 224 more rows
```

2.6 Grouped summaries with `summarise()`

Extract mean fuel economy for cities and highways:

```
mpg %>%
  summarise(mean_cty = mean(cty),
            mean_hwy = mean(hwy))
```

```
## # A tibble: 1 x 2
##   mean_cty mean_hwy
##      <dbl>    <dbl>
## 1     16.9     23.4
```

Extract mean fuel economy for cities and highways, by car `class`:

```
mpg %>%
  group_by(class) %>%
  summarise(mean_cty = mean(cty),
            mean_hwy = mean(hwy))
```

```
## # A tibble: 7 x 3
##   class      mean_cty mean_hwy
##   <chr>         <dbl>    <dbl>
## 1 2seater        15.4     24.8
## 2 compact        20.1     28.3
## 3 midsize        18.8     27.3
## 4 minivan        15.8     22.4
## 5 pickup         13       16.9
## 6 subcompact     20.4     28.1
## 7 suv            13.5     18.1
```

Note that we strung together two operations using the pipe. We can string together arbitrarily many operations using the pipe, including plotting:

```
mpg %>%
  group_by(class) %>%
  summarise(mean_cty = mean(cty),
            mean_hwy = mean(hwy)) %>%
  ggplot(aes(x = mean_cty, y = mean_hwy, label = class)) +
  geom_point() +
  ggrepel::geom_text_repel() +
  labs(x = "Mean city fuel economy (gallons per mile)",
```

```
      y = "Mean highway fuel economy (gallons per mile)") +
  theme_bw()
```

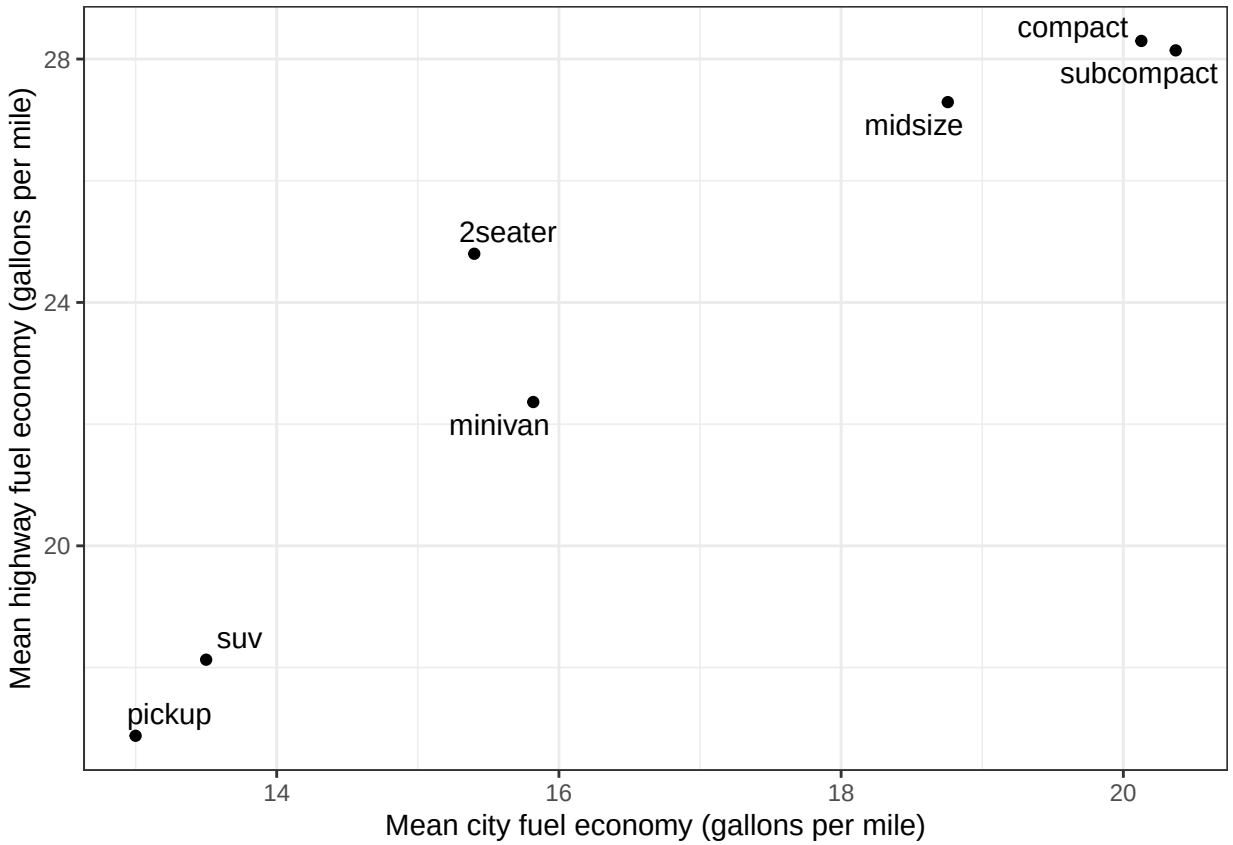

Common functions used with `summarise()`: `mean`, `median`, `sum`, `min`, `max`, `n`, `sd`, ...

3 Exploratory data analysis

All these visualizations and transformations can help us to explore the data and find interesting patterns.

3.1 Relationship between engine size and fuel economy

What is the relationship between `displ` (a car's engine size in liters) and `hwy` (a car's fuel efficiency on the highway, in miles per gallon)? We created the plot below already:

```
mpg %>%
  ggplot(aes(x = displ, y = hwy)) +
  geom_point(aes(color = class)) +
  geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

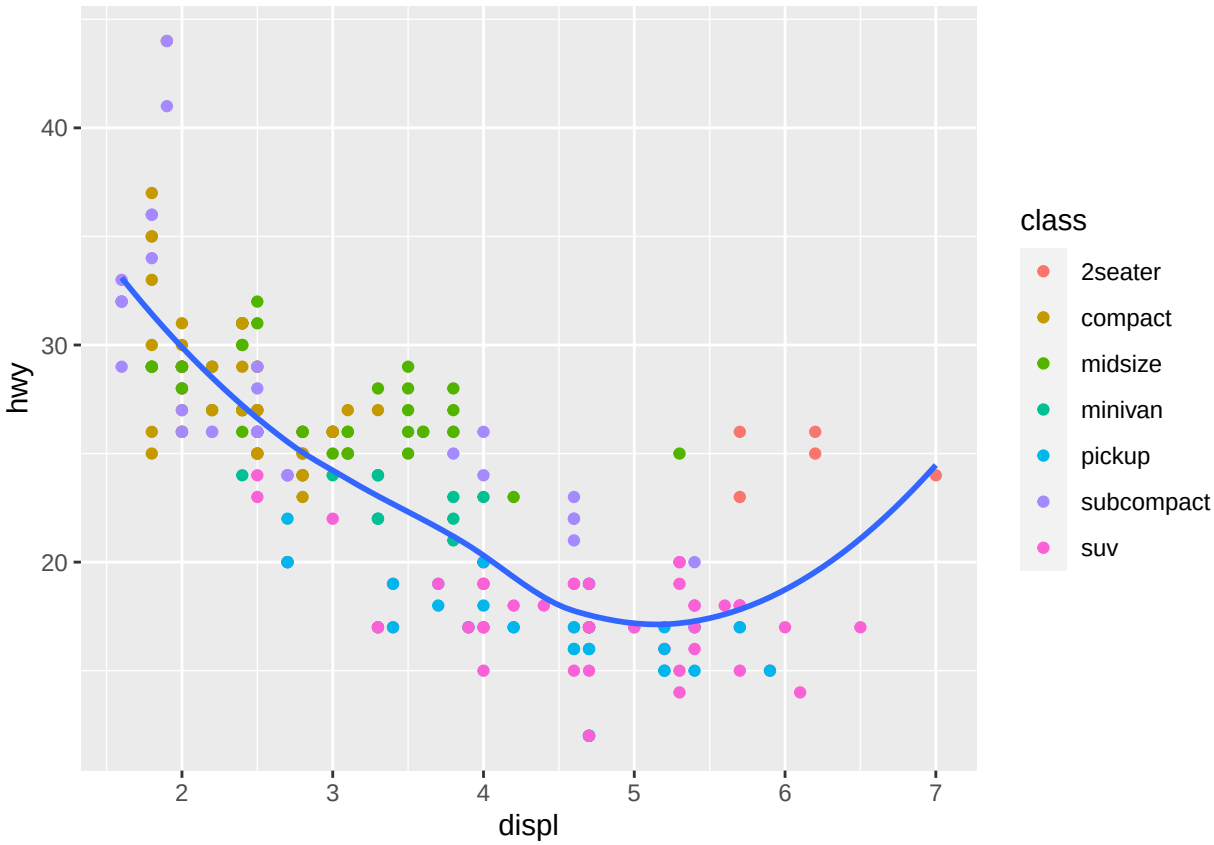

Note that the smooth curve is first decreasing but then increasing. What would we expect the relationship to be between `hwy` and `displ`? What color points seem to be "pulling up" the smooth curve fit? Why might this be the case?

3.2 Fuel economy in cities versus on highways

What is the relationship between fuel economy in cities and on highways?

3.3 Comparing manufacturers based on fuel economy

Which manufacturers had the best and worst highway fuel economy in the year 1999, on average over models? What about the year 2008?