

# Unit 3 Lecture 1: Logistic Regression

October 5, 2021

```
library(pROC)           # for ROC curves
library(tidyverse)
```

In today's R demo, we will apply logistic regression to the `Default` data from lecture:

```
default_data = ISLR2::Default %>% as_tibble()
default_data
```

```
## # A tibble: 10,000 x 4
##   default student balance income
##   <fct>   <fct>     <dbl>  <dbl>
## 1 No      No         730.  44362.
## 2 No      Yes         817.  12106.
## 3 No      No        1074.  31767.
## 4 No      No         529.  35704.
## 5 No      No         786.  38463.
## 6 No      Yes         920.   7492.
## 7 No      No         826.  24905.
## 8 No      Yes         809.  17600.
## 9 No      No        1161.  37469.
## 10 No     No           0   29275.
## # ... with 9,990 more rows
```

As an exploratory question, what is the default rate in this data?

The rest of the activity will be easier if we code `default` as 0-1:

```
default_data = default_data %>% mutate(default = as.numeric(default == "Yes"))
default_data
```

```
## # A tibble: 10,000 x 4
##   default student balance income
##   <dbl> <fct>     <dbl>  <dbl>
## 1      0 No         730.  44362.
## 2      0 Yes         817.  12106.
## 3      0 No        1074.  31767.
## 4      0 No         529.  35704.
## 5      0 No         786.  38463.
## 6      0 Yes         920.   7492.
## 7      0 No         826.  24905.
## 8      0 Yes         809.  17600.
## 9      0 No        1161.  37469.
## 10     0 No           0   29275.
## # ... with 9,990 more rows
```

Let's split the default data into training and test sets:

```
set.seed(471)
train_samples = sample(1:nrow(default_data), 0.8*nrow(default_data))
default_train = default_data %>% filter(row_number() %in% train_samples)
default_test = default_data %>% filter(!(row_number() %in% train_samples))
```

## Running a logistic regression

The way to run a logistic regression is through the `glm` function:

```
glm_fit = glm(default ~ student + balance + income,
              family = "binomial",
              data = default_train)
summary(glm_fit)
```

```
##
## Call:
## glm(formula = default ~ student + balance + income, family = "binomial",
##      data = default_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1433  -0.1416  -0.0541  -0.0195   3.7396
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.106e+01  5.537e-01 -19.972  < 2e-16 ***
## studentYes  -7.051e-01  2.663e-01  -2.648  0.00811 **
## balance      5.824e-03  2.595e-04  22.444  < 2e-16 ***
## income       5.806e-06  9.112e-06   0.637  0.52399
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2394.2  on 7999  degrees of freedom
## Residual deviance: 1259.9  on 7996  degrees of freedom
## AIC: 1267.9
##
## Number of Fisher Scoring iterations: 8
```

## Interpreting the estimates

- What is the coefficient estimate for `student`?
- Does this suggest that being a student increases or decreases the probability of default, other things being equal?
- According to this estimate, how does being a student impact the log-odds of default? How does it impact the odds of default?

## Extracting elements of the fit

We can extract the coefficient estimates, standard errors, etc. just as we did with linear models:

```
coef(glm_fit)
```

```
## (Intercept) studentYes balance income
## -1.105920e+01 -7.050610e-01 5.824362e-03 5.806162e-06
```

## Fitted probabilities and making predictions

We can extract the fitted probabilities of default for a test set using the `predict` function:

```
fitted_probabilities = predict(glm_fit,
  newdata = default_test,
  type = "response")      # to get output on probability scale
head(fitted_probabilities)
```

```
##      1      2      3      4      5      6
## 8.830647e-06 1.167429e-02 1.082938e-04 1.860563e-04 8.007356e-05 7.896278e-04
```

We can now make predictions based on the fitted probabilities using the standard 0.5 threshold:

```
predictions = as.numeric(fitted_probabilities > 0.5)
head(predictions)
```

```
## [1] 0 0 0 0 0 0
```

## Evaluating the classifier

Let's calculate the misclassification rate of the above logistic regression classifier.

```
# first add predictions to the tibble
default_test = default_test %>%
  mutate(predicted_default = predictions)
default_test
```

```
## # A tibble: 2,000 x 5
##   default student balance income predicted_default
##   <dbl> <fct>      <dbl> <dbl>      <dbl>
## 1      0 Yes         0  21871.         0
## 2      0 No       1113.  23810.         0
## 3      0 No        286.  45042.         0
## 4      0 Yes        528.  17637.         0
## 5      0 No        229.  50500.         0
## 6      0 No        642.  30466.         0
## 7      0 No        773.  34353.         0
## 8      0 Yes        221.  16873.         0
## 9      0 No        409.  54207.         0
## 10     0 No       1228.  37409.         0
## # ... with 1,990 more rows
```

```
# then calculate misclassification rate
default_test %>%
  summarise(mean(default != predicted_default))
```

```
## # A tibble: 1 x 1
##   `mean(default != predicted_default)`
##   <dbl>
## 1      0.029
```

To get a fuller picture, let's calculate the confusion matrix:

```
default_test %>%
  select(default, predicted_default) %>%
  table()
```


```
##           predicted_default
## default    0      1
##           0 1931   11
##           1   47   11
```

- What are the false positive and false negative rates of this classifier?
- If the cost of a false negative is three times that of a false positive, what probability threshold should we use? What are the false positive and false negative rates for the resulting classifier?

Next, let's plot the ROC curve for this classifier.

```
# ROC curve
roc_data = roc(default_test %>% pull(default),
               fitted_probabilities)

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
tibble(FPR = 1-roc_data$specificities,
       TPR = roc_data$sensitivities) %>%
  ggplot(aes(x = FPR, y = TPR)) +
  geom_line() +
  geom_abline(slope = 1, linetype = "dashed") +
  # geom_point(x = fpr, y = 1-fnr, colour = "red") +
  theme_bw()
```




```
# print the AUC
roc_data$auc
```

```
## Area under the curve: 0.9438
```

## Plotting a univariate logistic regression fit

Univariate logistic regression fits can be plotted using `geom_smooth`:

```
default_train %>%
  ggplot(aes(x = balance, y = default)) +
  geom_jitter(height = .05) +
  geom_smooth(method = "glm",
              formula = "y~x",
              method.args = list(family = "binomial"),
              se = FALSE) +
  ylab("Prob(default=1)") +
  theme_bw()
```



Roughly at what value of balance do we switch from predicting no default to predicting default?